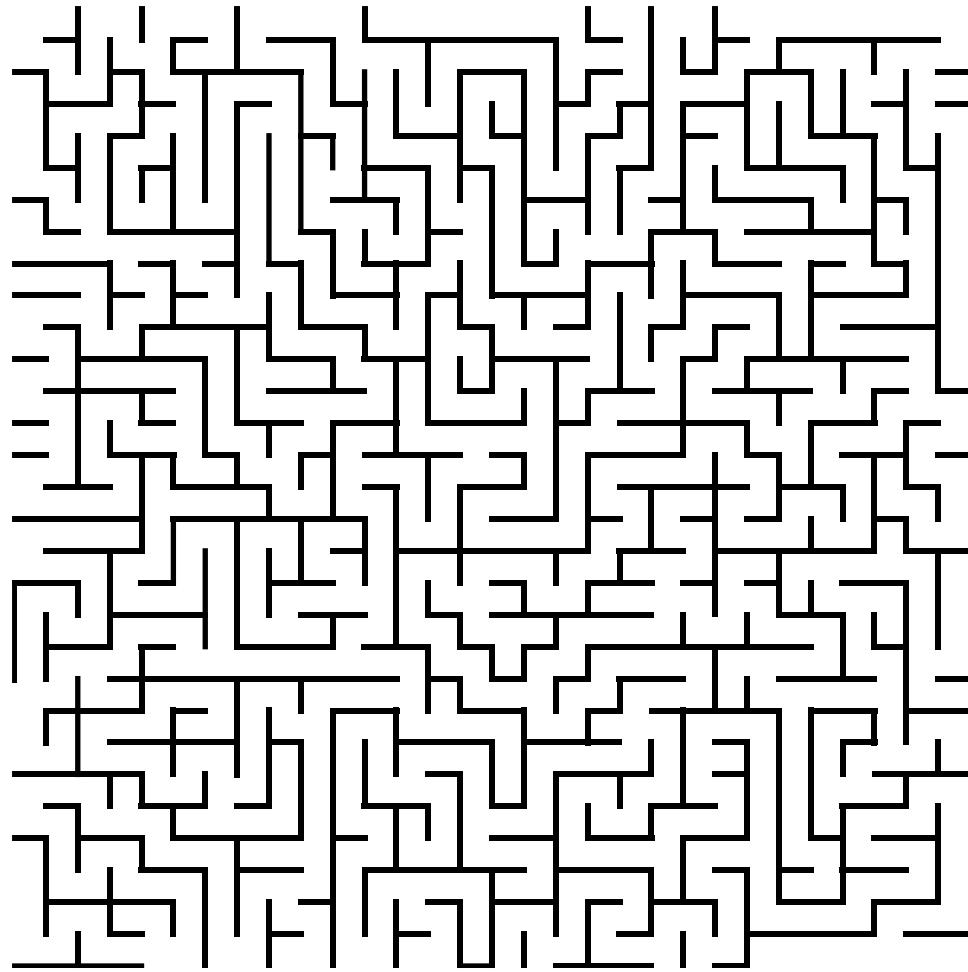




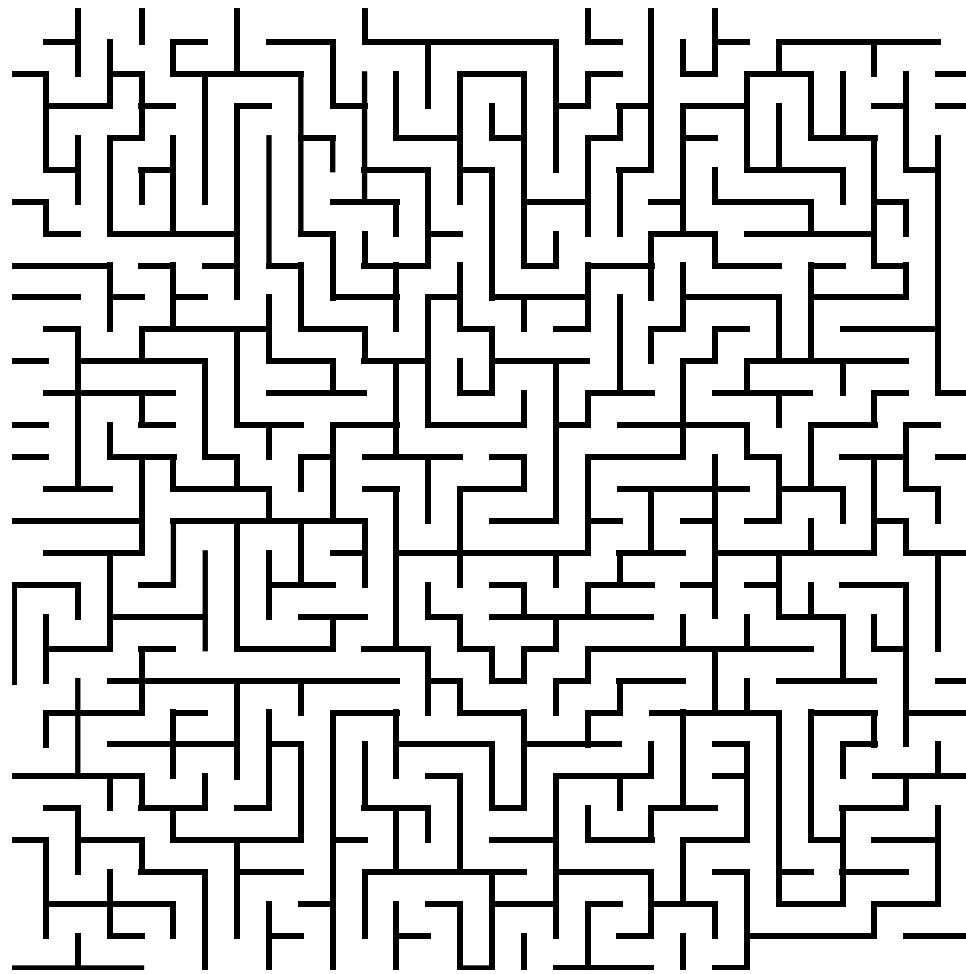
# Generative AI

2025 / 10 / 15



Prompting is a collaborative effort between the LLM and you to navigate through the maze of possibilities.

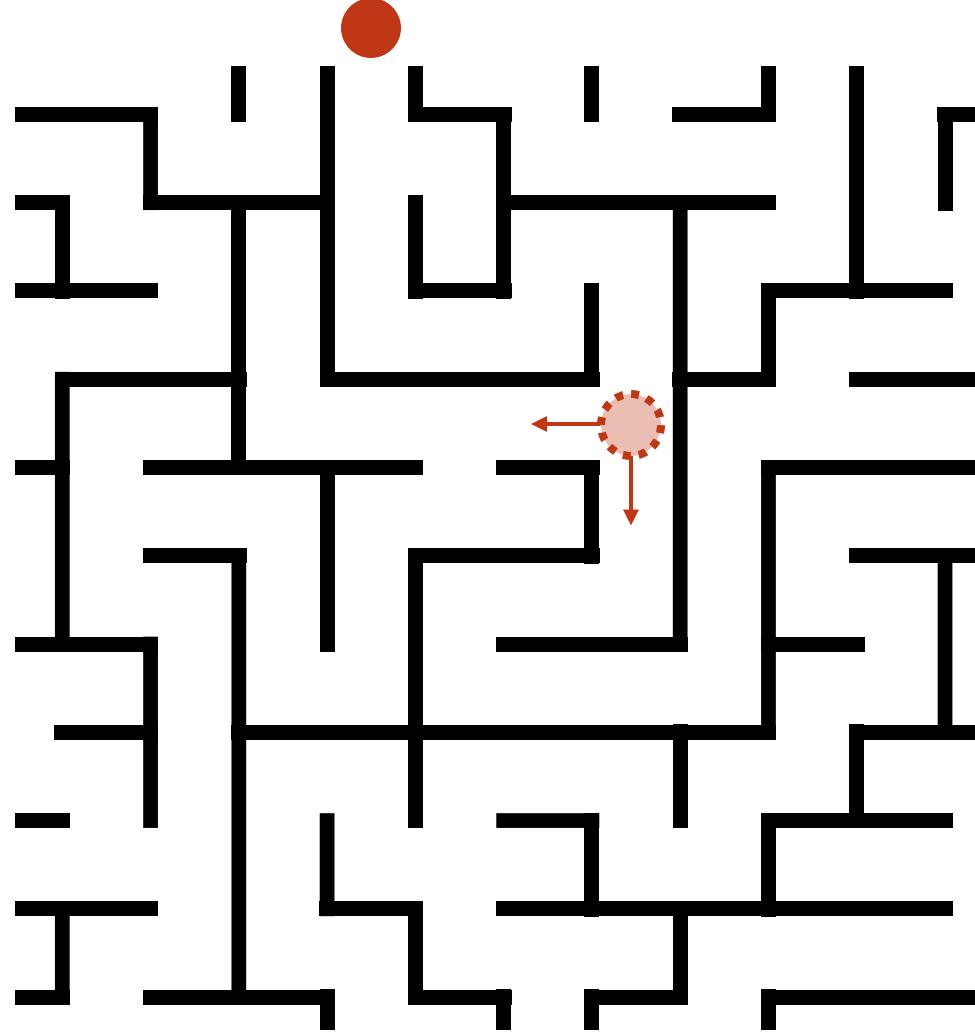
Anything you add to the prompt will explicitly affect the output. Anything you don't add to the prompt will implicitly affect the output.<sup>1</sup>

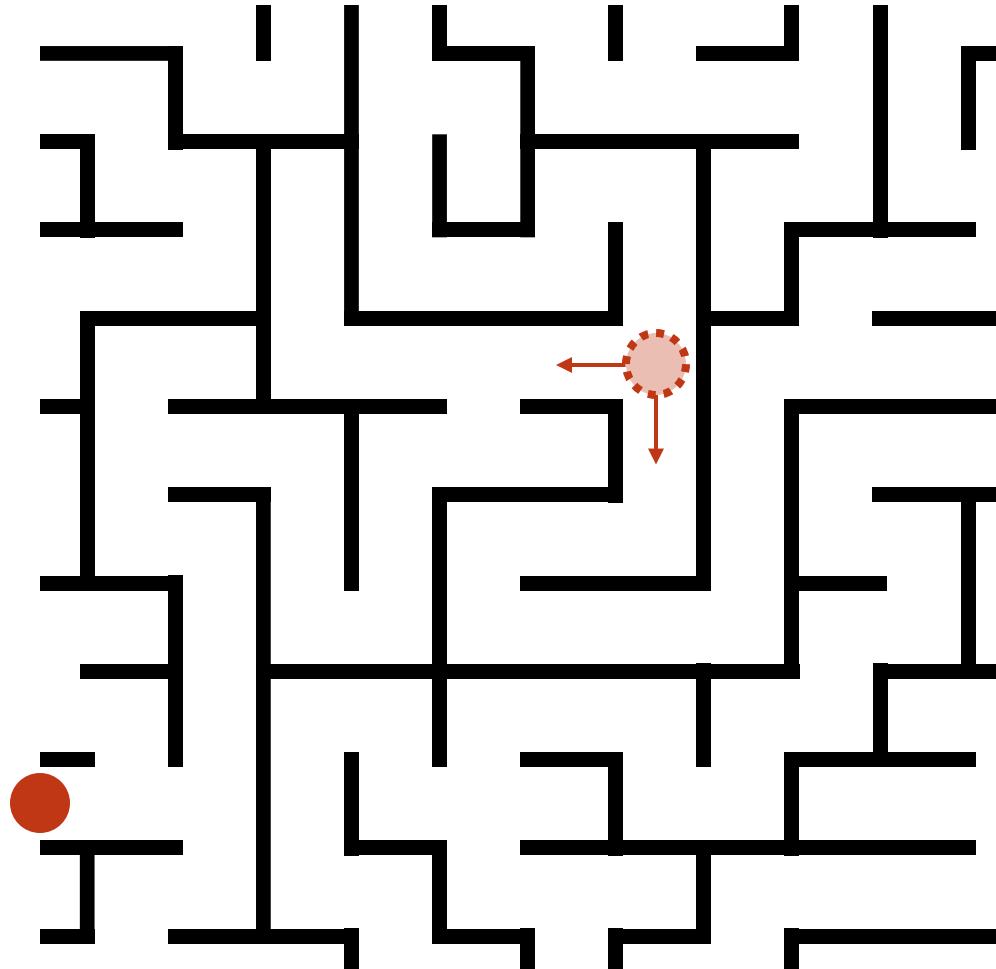


Write me a function to sort a list

+ ↗

OLMo 2 ↕ ↑





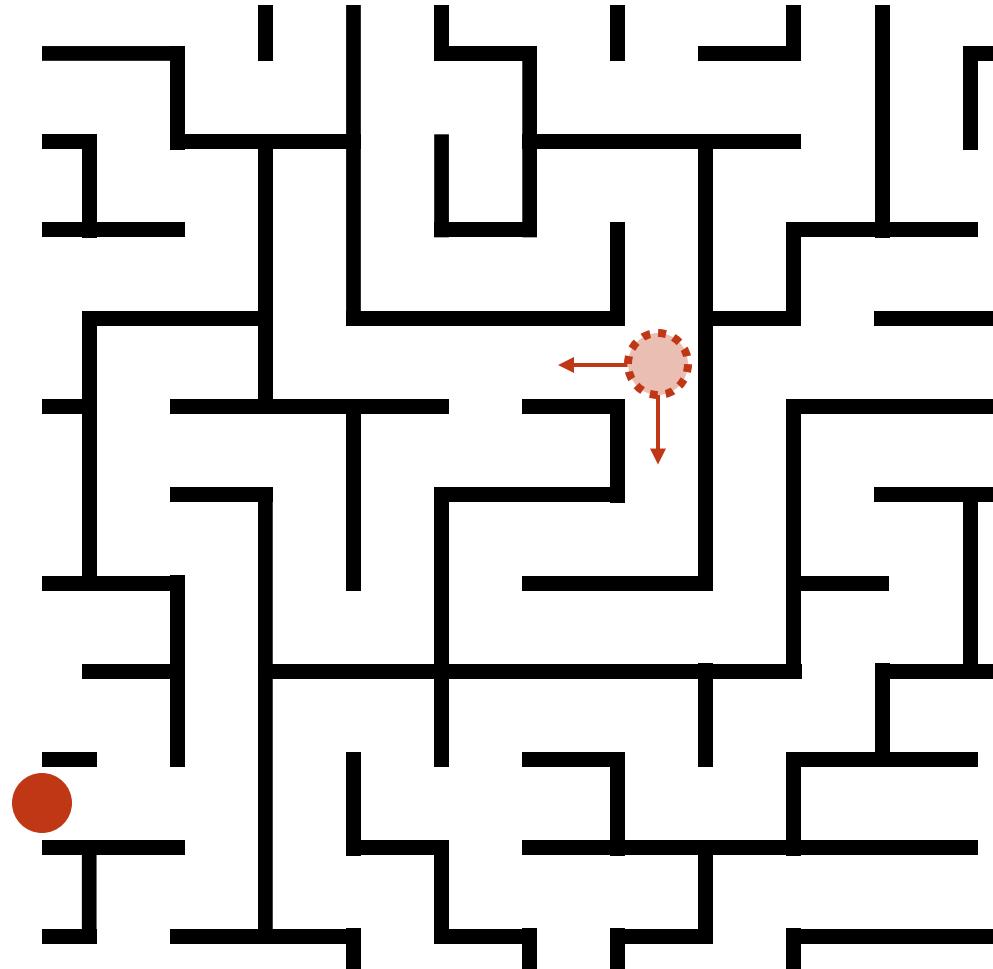
## Quick Fire Prompting Tipps



Let the LLM explain its decisions.

### Examples

- ... What implicit decisions did you make?
- ... What would I need to change in my prompt to achieve  $XYZ$ .



## Quick Fire Prompting Tipps



Let the LLM explain its decisions.



Provide more **context**

- Your mental model / perspective / assumptions
- Previous solutions you tried
- Hard constraints and soft preferences

# AGENDA

1

**What** is (non-)generative AI?

2

A bird's eye view of **generative sequence modeling**.

3

A bird's eye view of **diffusion modeling**.

4

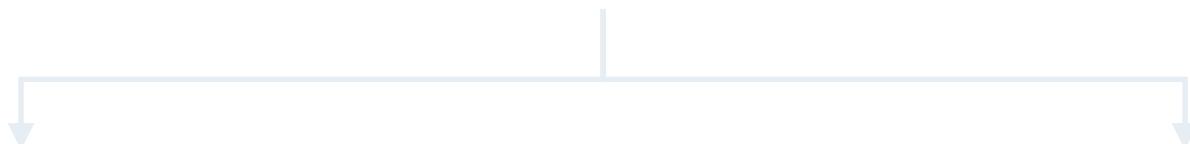
A bird's eye view of the **transformer** architecture.

What is  
*Generative Artificial Intelligence*

## **Artificial Intelligence**

**Symbolic AI**

**Subsymbolic AI**





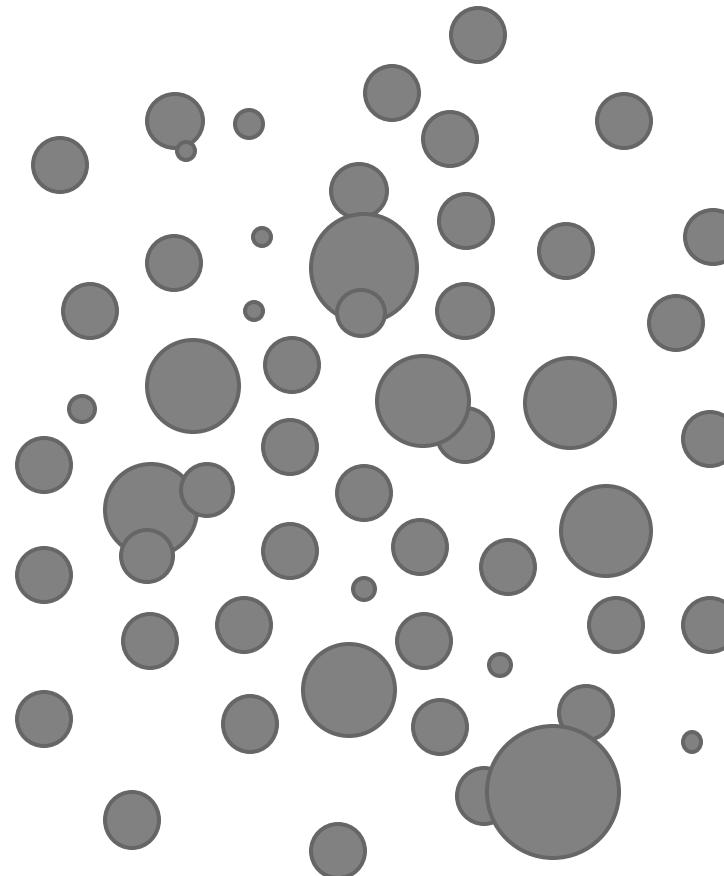


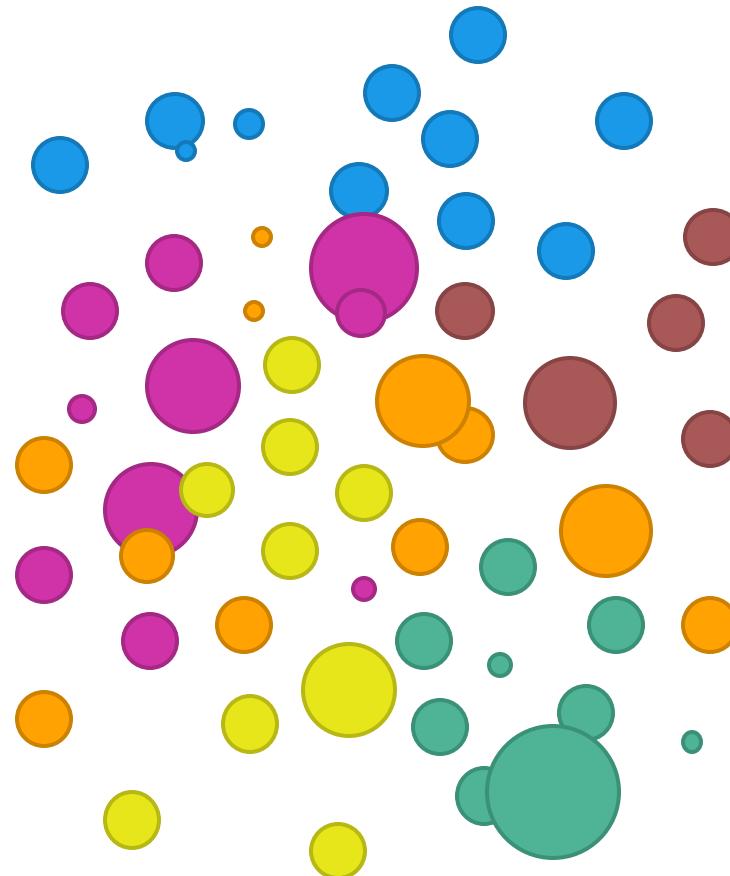
*“A book is either available or a member has borrowed it.”*

$\forall x (\text{Book}(x) \rightarrow (\text{Available}(x) \vee \exists y (\text{Member}(y) \wedge \text{Borrowed}(y, x))))$

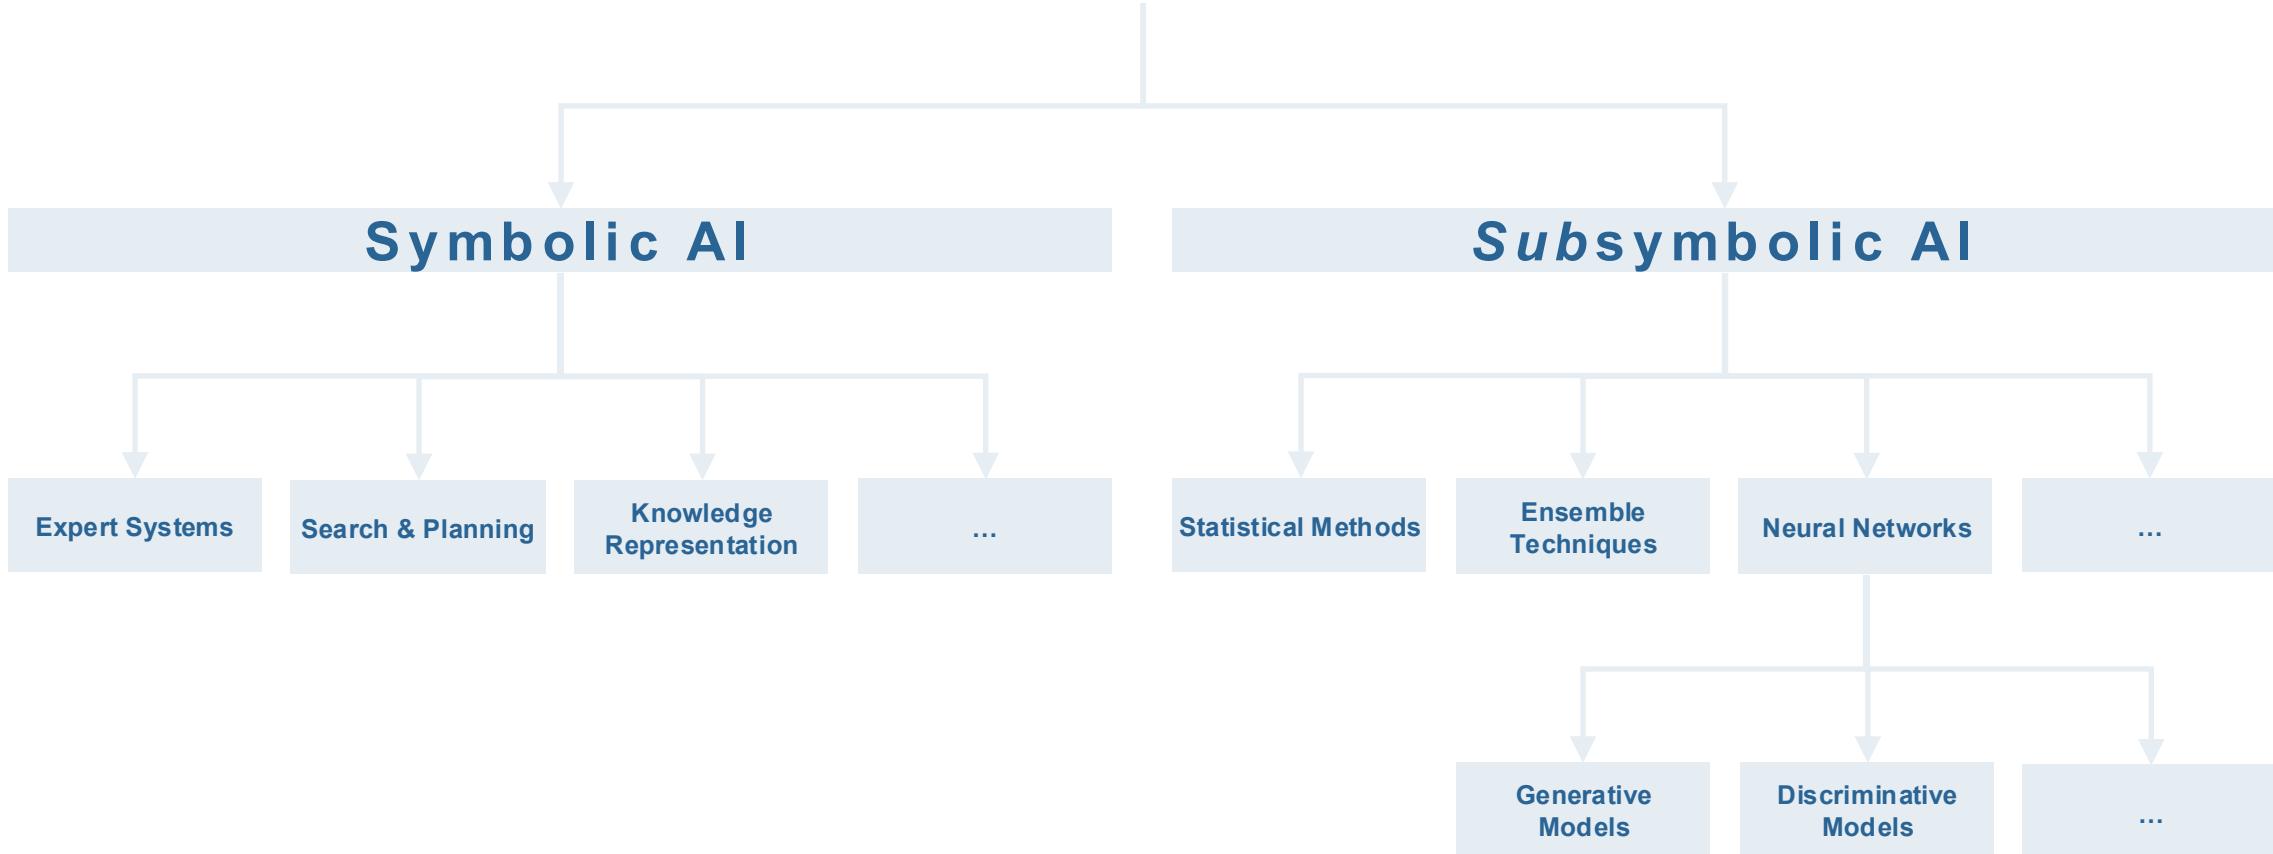
*“Is a book x available?”*

$\neg \exists y (\text{Member}(y) \wedge \text{Borrowed}(y, x))$

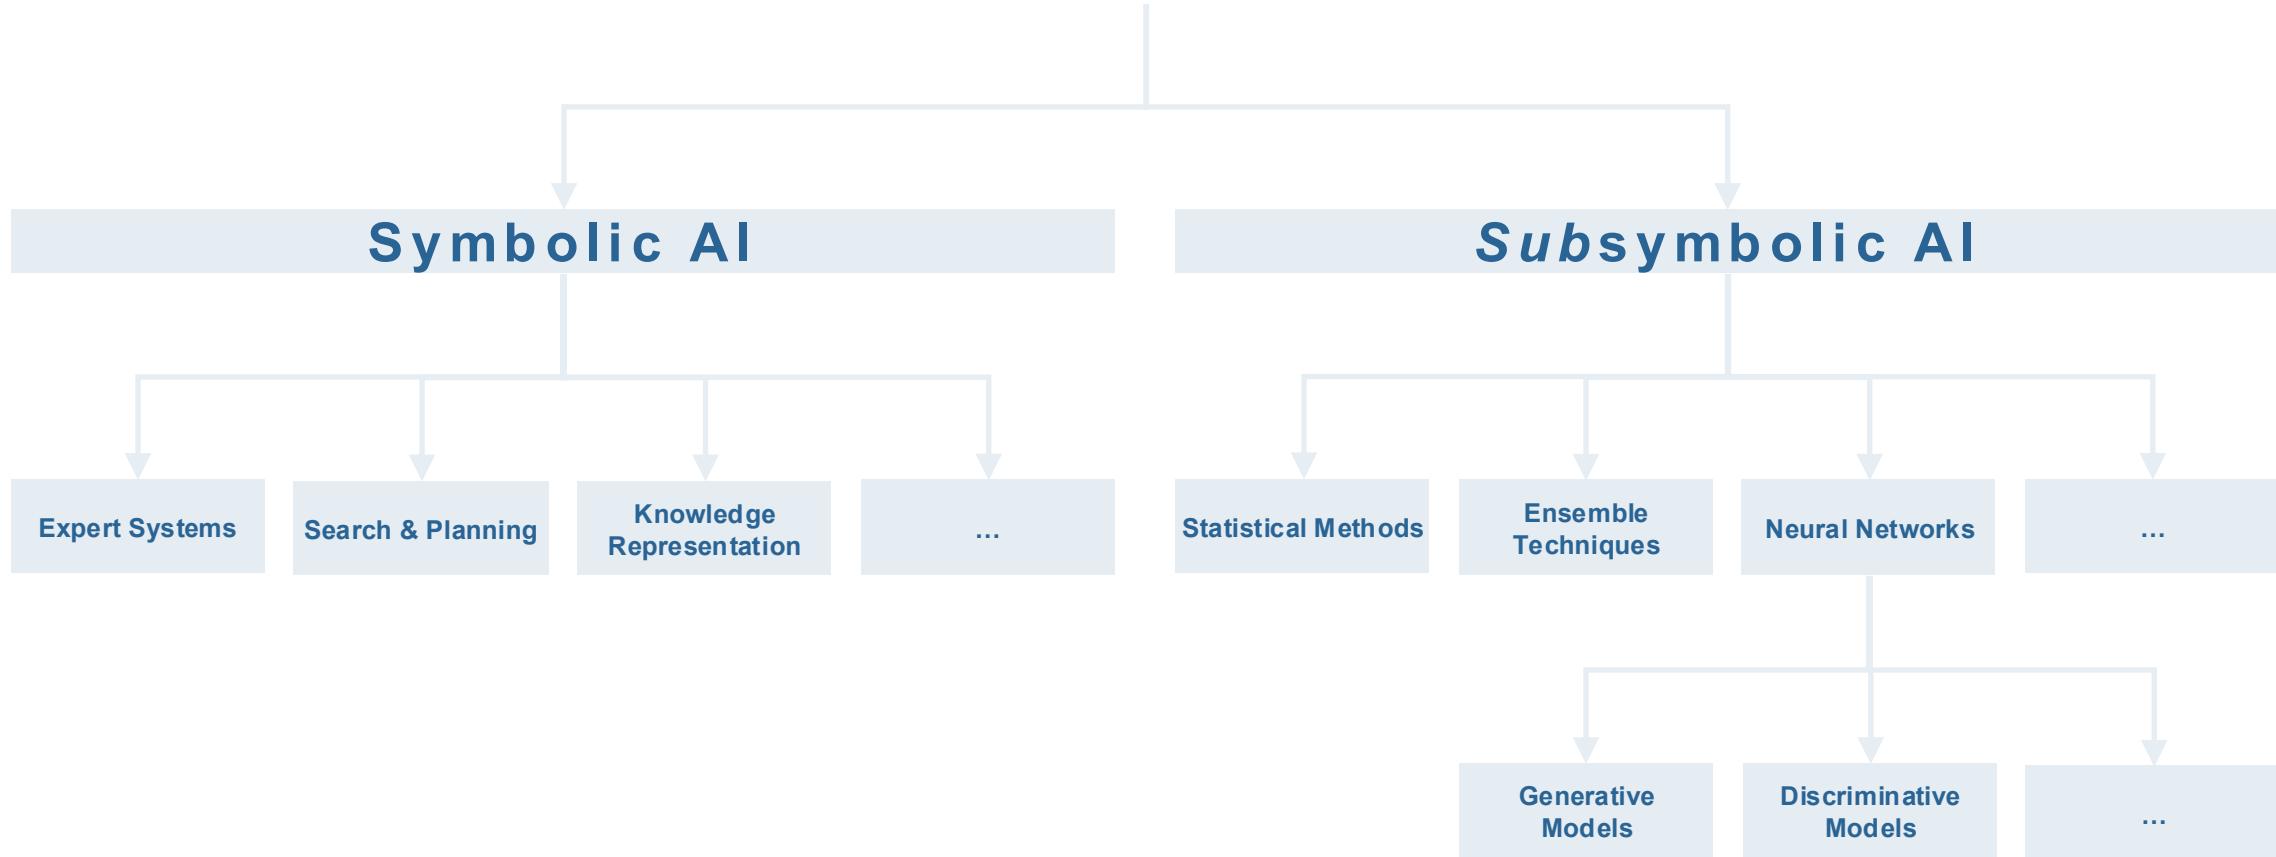


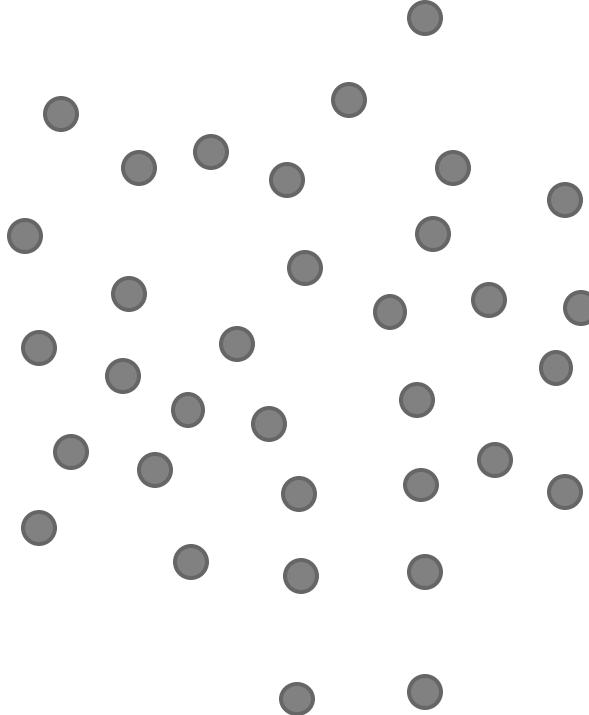


# Artificial Intelligence



# Artificial Intelligence

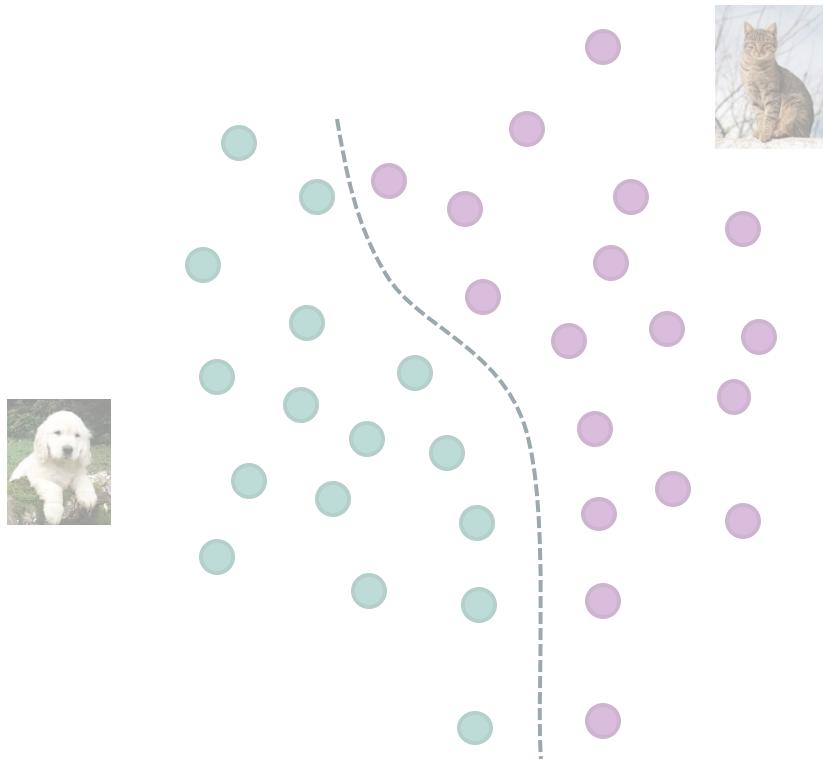






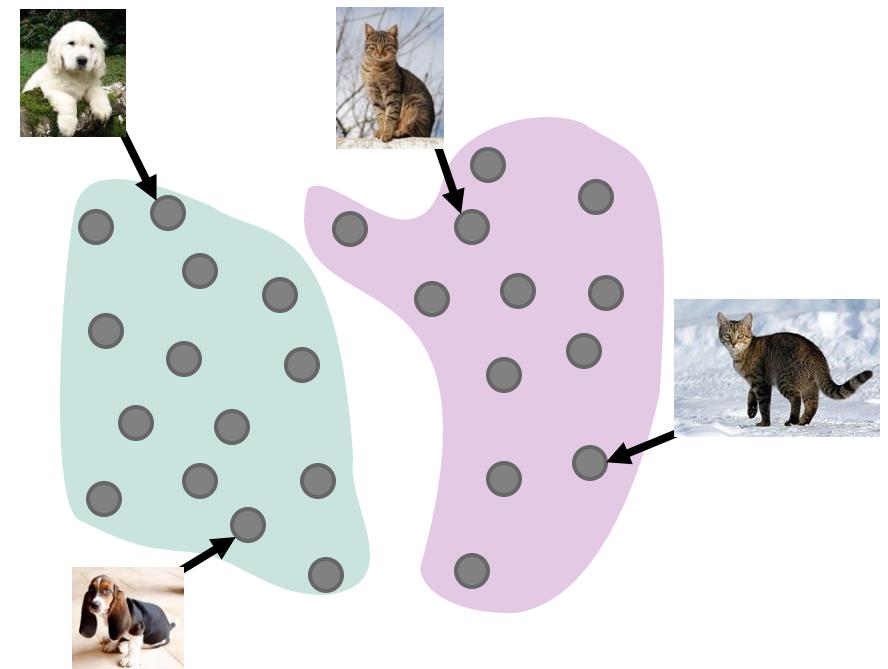
Discriminative Models

$$p(Y|X)$$



Discriminative Models

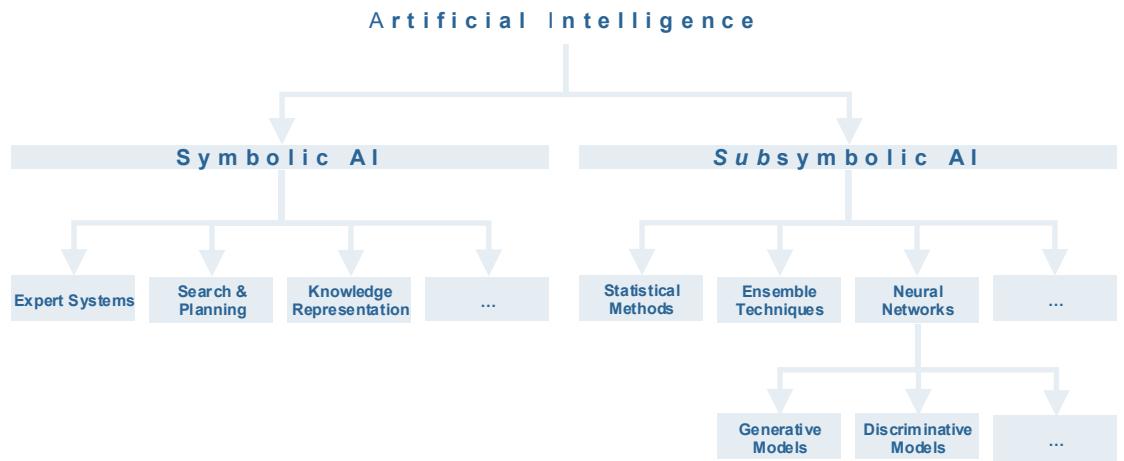
$$p(Y|X)$$



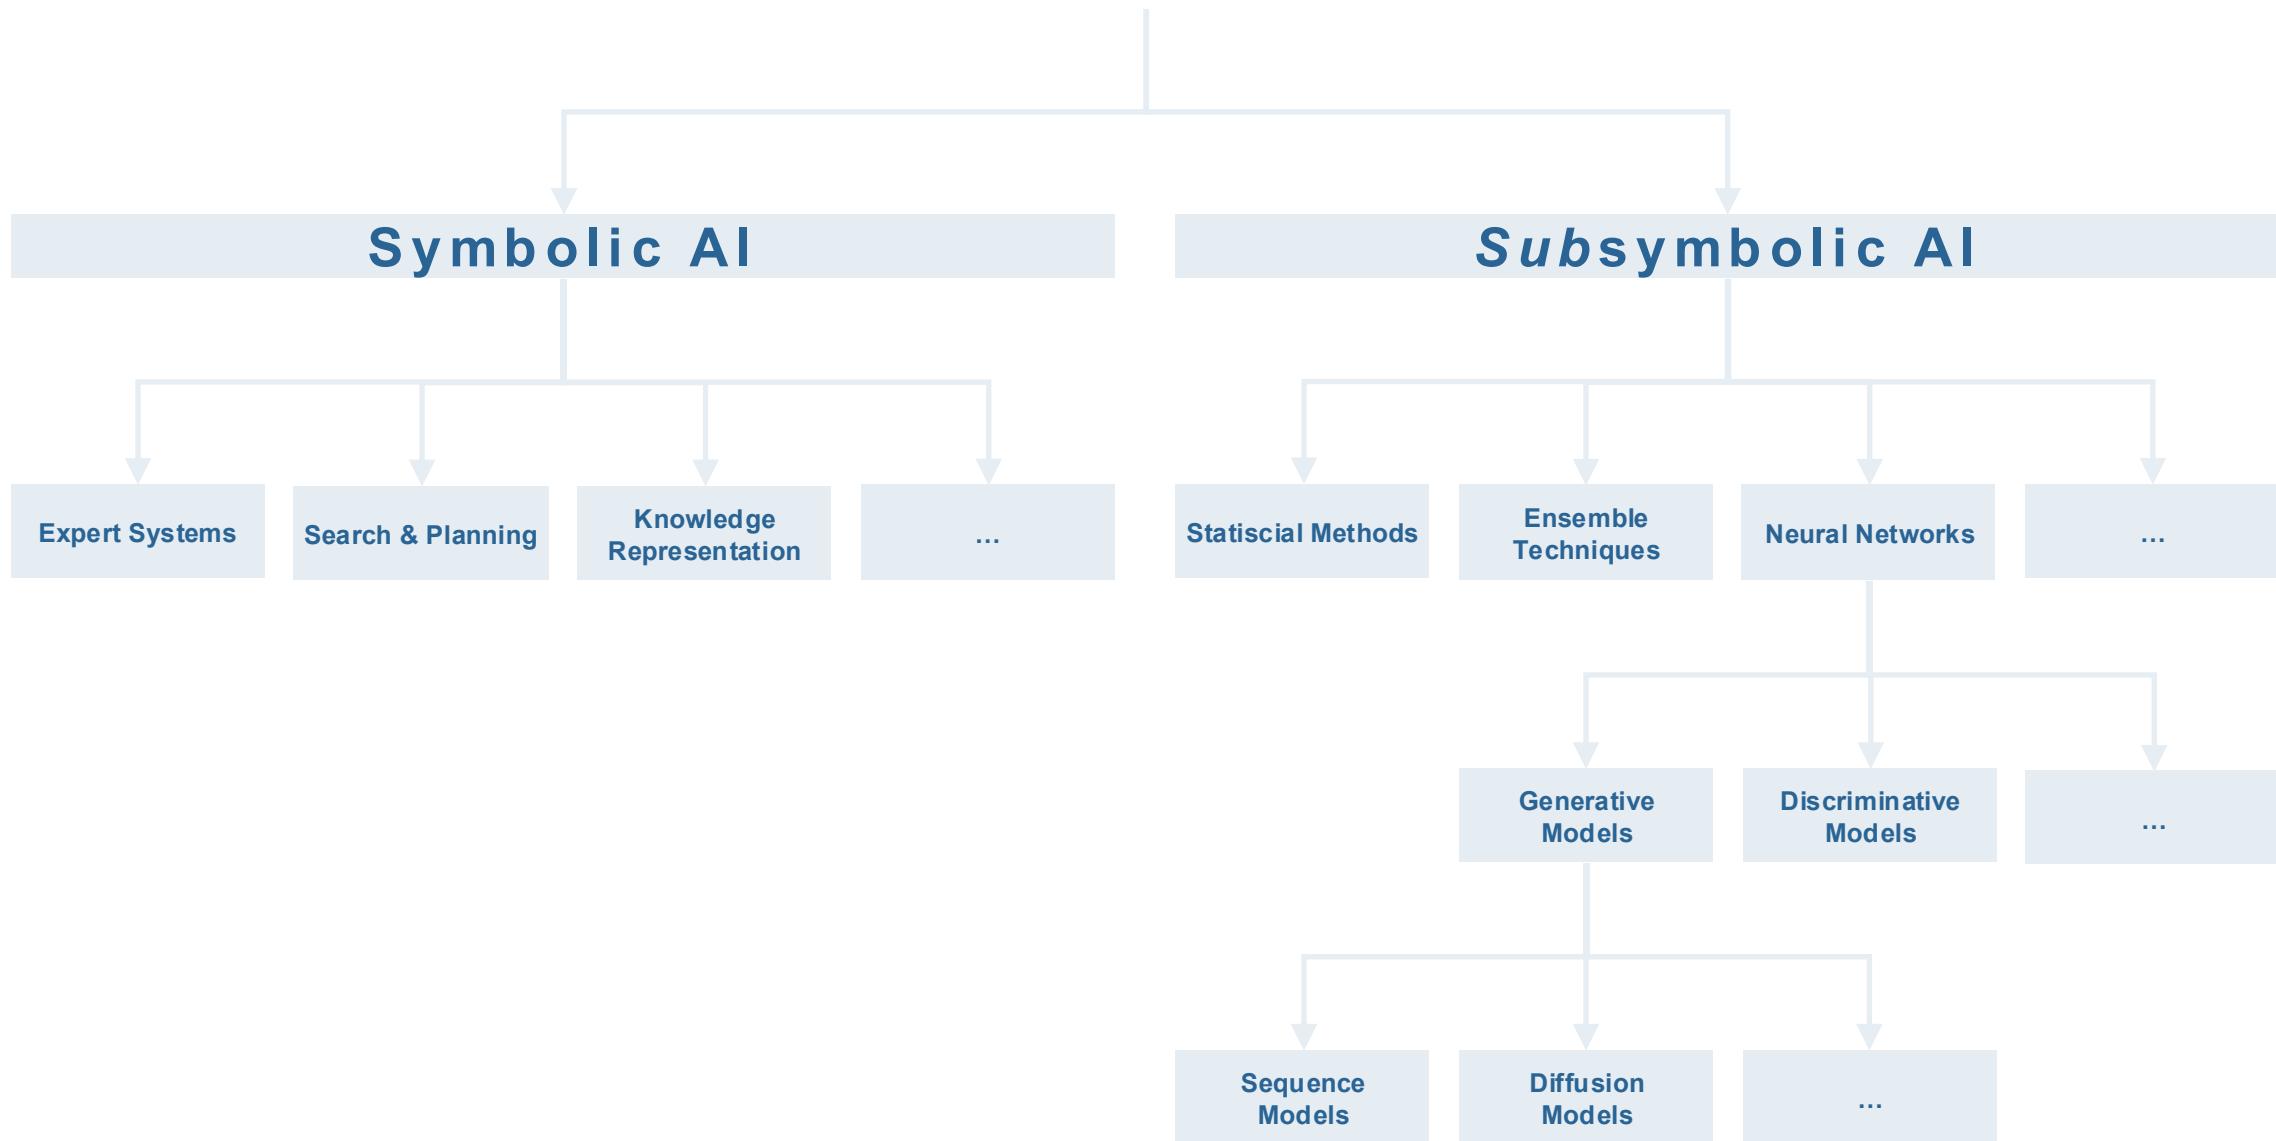
Generative Models

$$p(X,Y) \text{ or } p(X)$$

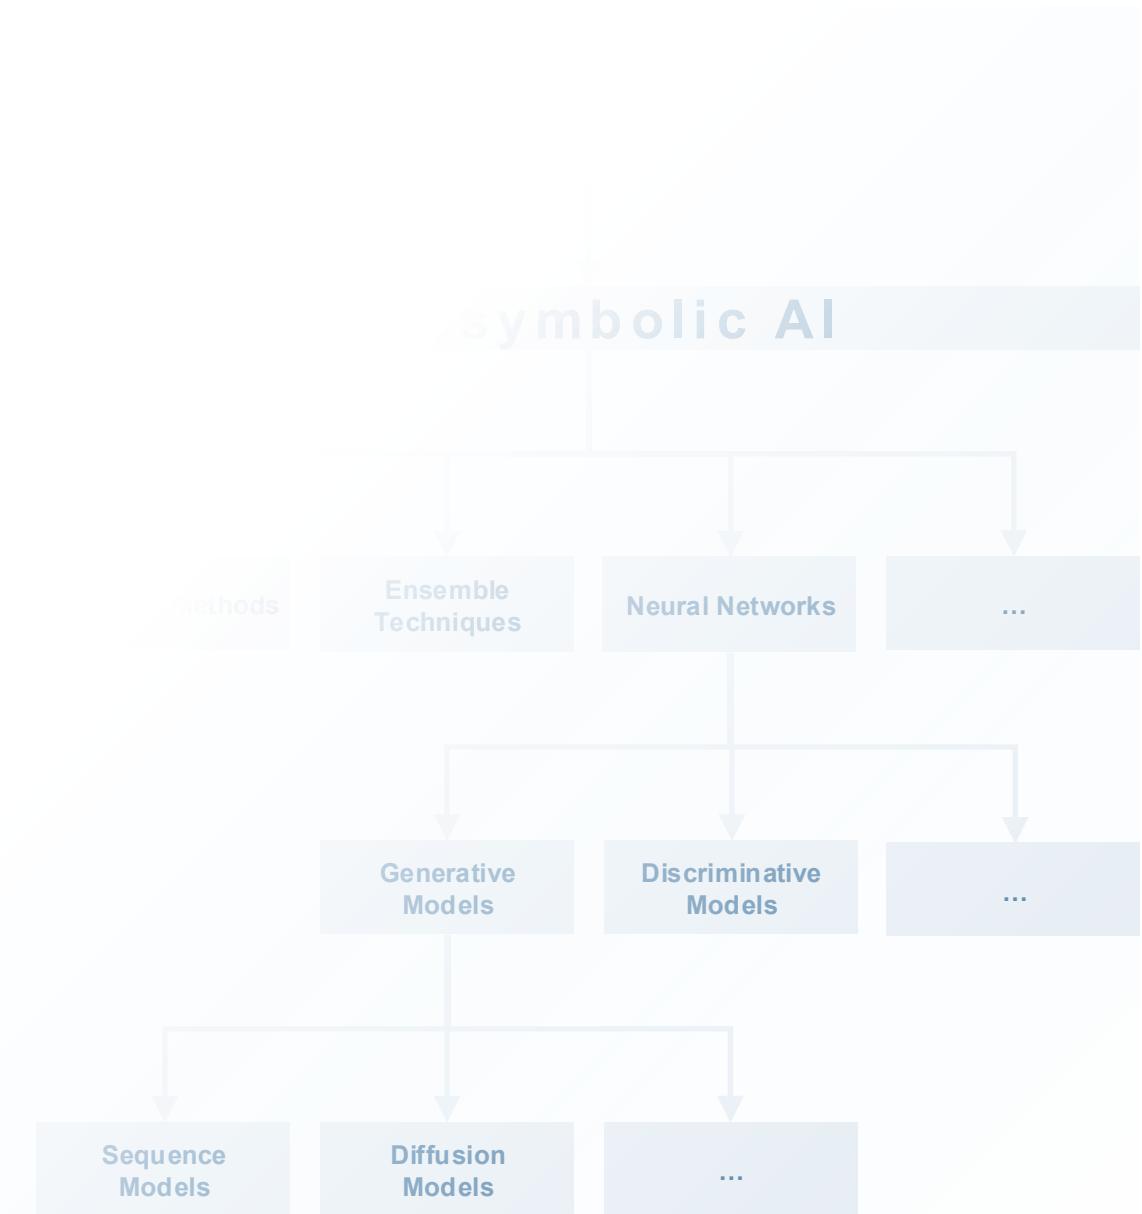
# What is (non-)generative AI? Recap



# Artificial Intelligence



# Sequence Modeling Tokens



The cute green dragon trotted into the cave

The cute green dragon trot ted into the cave

Token ID

72 302 902 1041 78 15 982 5 873



Token

The cute green dragon trot ted into the cave



The goal of generative language modelling is to capture the “true distribution of language”:  $p(X)$

### A few definitions:

- $V$  ... vocabulary of unique tokens  $s_i \in V$
- $n$  ... maximum sequence length
- $X$  ... space of all possible token sequences

$X = \{(s_1, s_2, \dots, s_k) | s_i \in V, 0 \leq k \leq n, k \in \mathbb{Z}\}$  where  $k$  is the sequence length

- $x \in X$  ... a particular sequence of tokens
- The probability of a given sequence  $x$

$$\begin{aligned} p(x) &= p(s_1) \cdot p(s_2 | s_1) \cdot p(s_3 | s_1, s_2) \cdot \dots \cdot p(s_k | s_1, \dots, s_{k-1}) \\ &= p(s_1) \cdot \prod_{i=2}^k p(s_i | s_1, \dots, s_{i-1}) \end{aligned}$$

The probability of a sequence is the product of the probability of the first token and the conditional probabilities of each subsequent token given all previous tokens.

 The probability of a sequence is the product of the probability of the first token and the conditional probabilities of each subsequent token given all previous tokens.

### A few more definitions:

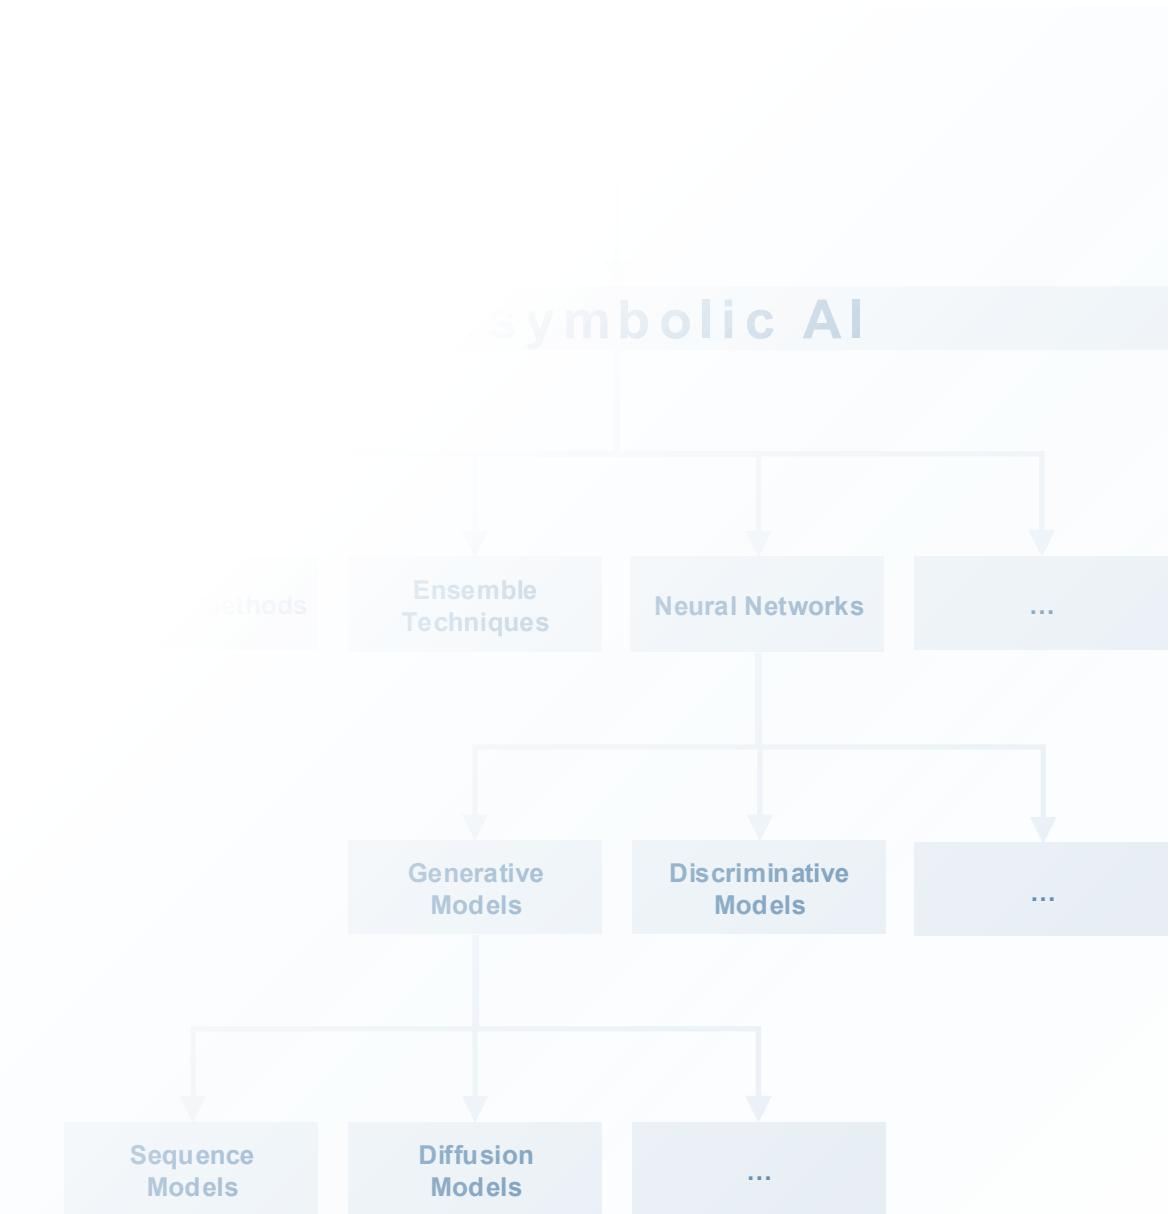
- $V_i$  ... the random variable for a token at position  $i$ , and  $s_i$  is the actual token observed at that position.
- $p(V_i | s_1, \dots, s_{i-1})$  ... the probability distribution over all possible tokens that will come next.
- $y_i$  ... a one-hot encoded vector of dimension equal to the vocabulary size holding 1 for the correct token at position  $i$  and 0 otherwise. This acts as the ground truth.

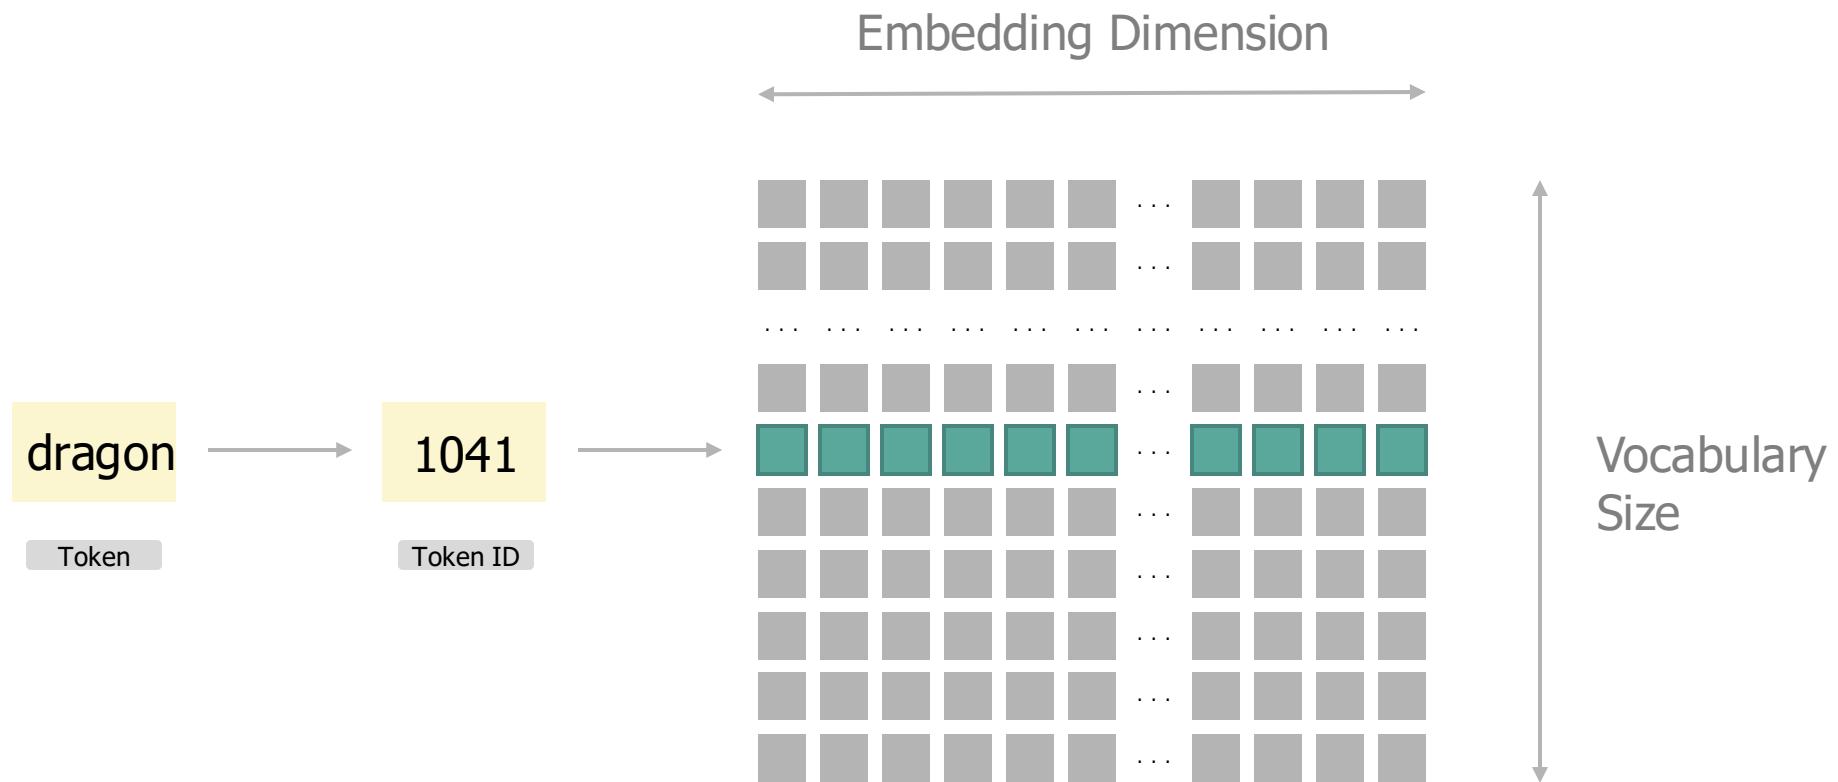
 
$$L = -\sum_i y_i \log(p(V_i | s_1, \dots, s_{i-1})) = -\sum_i \log(p(s_i | s_1, \dots, s_{i-1}))$$

 The loss function is defined as the negative log-likelihood of the correct tokens across all positions in the sequence, given the preceding tokens.

# Tokens in Sequence Modeling Recap

# Sequence Modelling Embeddings





The

cute

green

dragon

trot

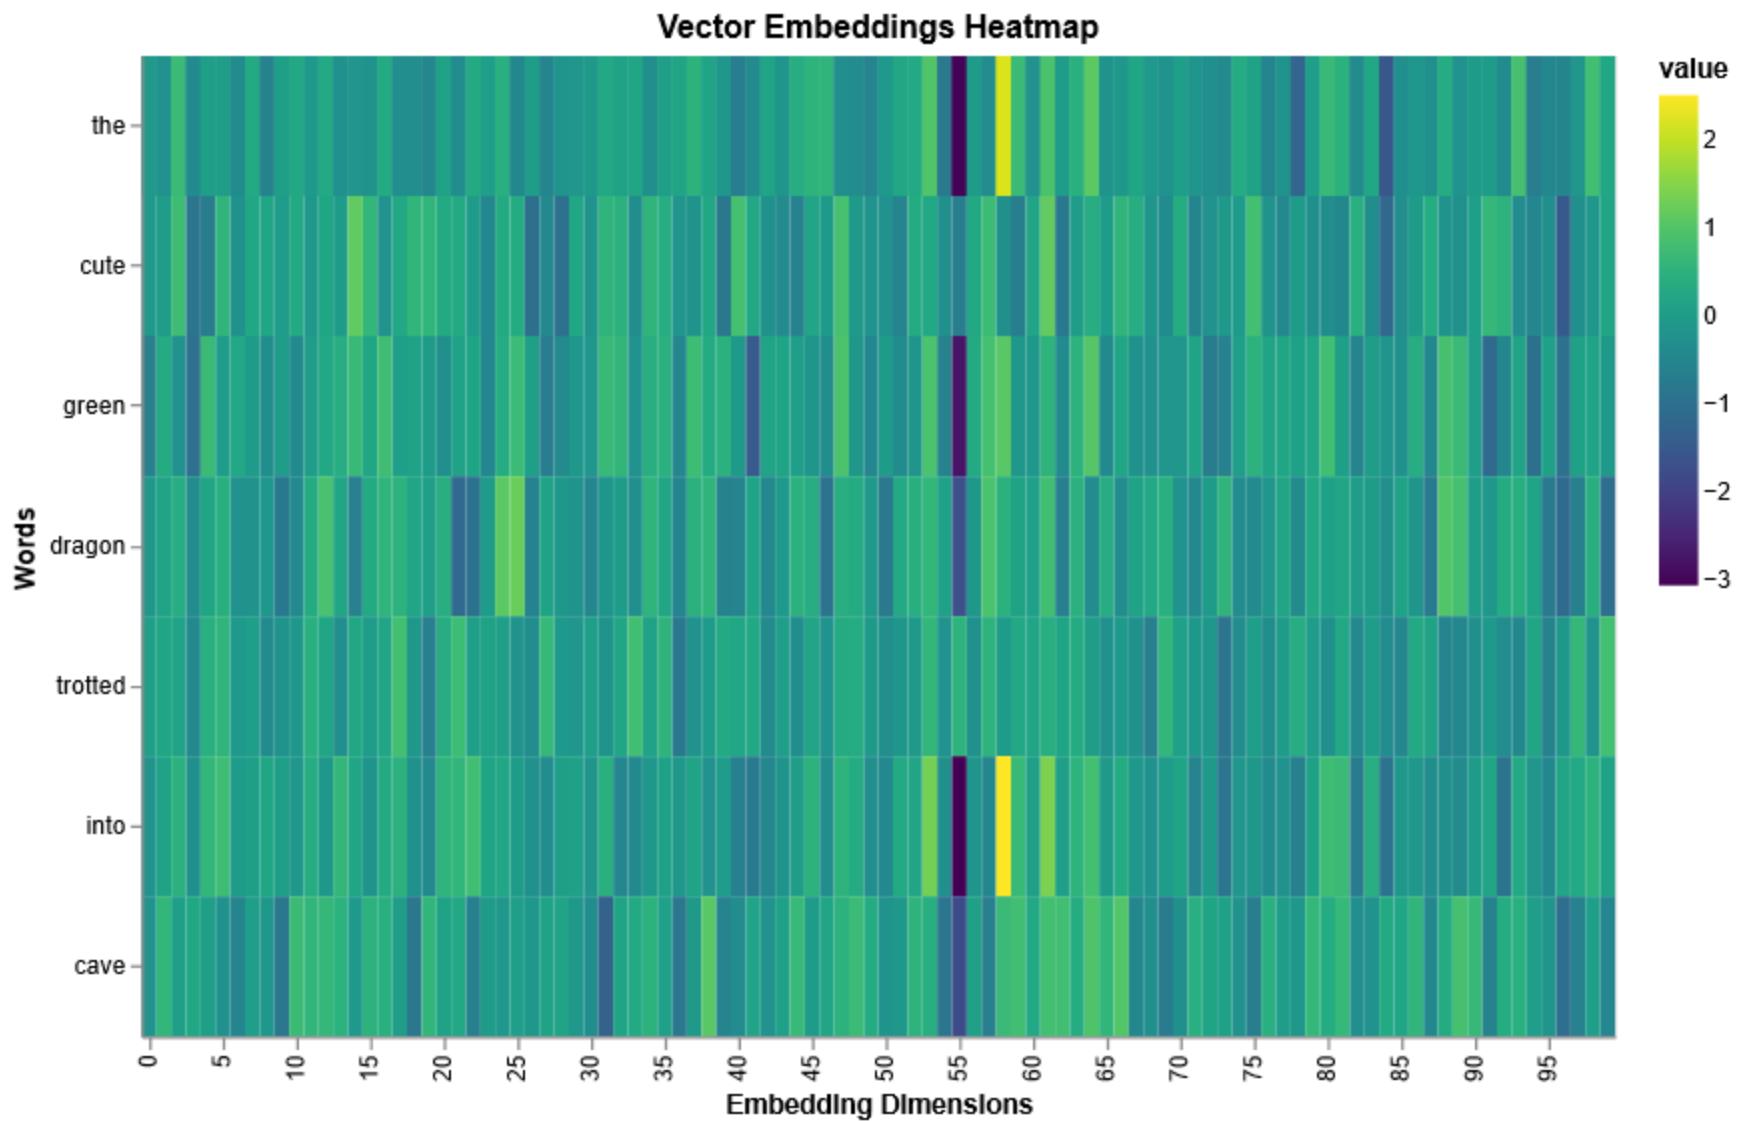
0.8
5.4
1.2
0.9
0.2
1.1
7.3
4.4
8.9
0.4
1.5
2.2
:
0.3

1.3
4.1
2.0
0.3
1.3
7.9
0.3
0.1
0.1
0.8
5.3
1.1
:
8.8

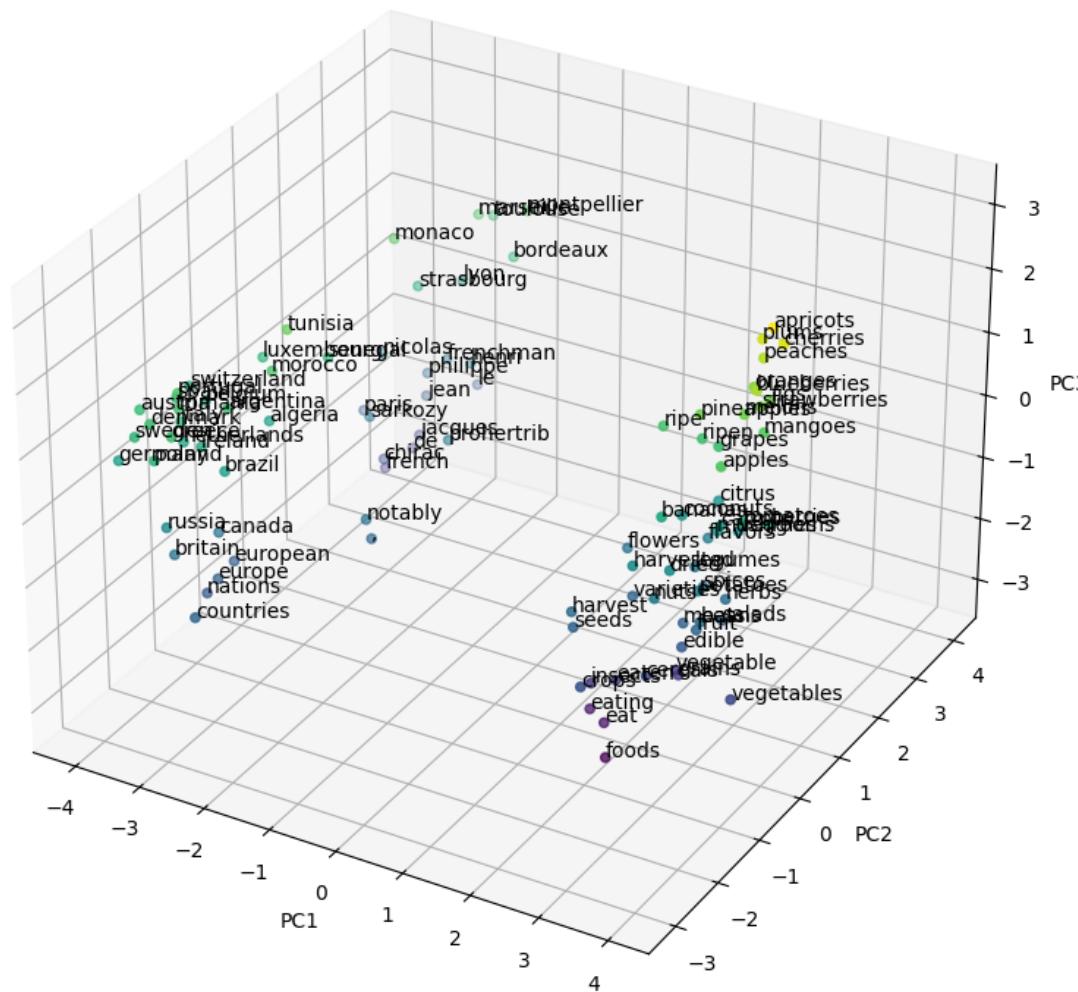
4.5
9.9
1.2
1.1
1.9
9.9
2.6
4.4
7.0
0.3
3.3
3.3
:
3.2

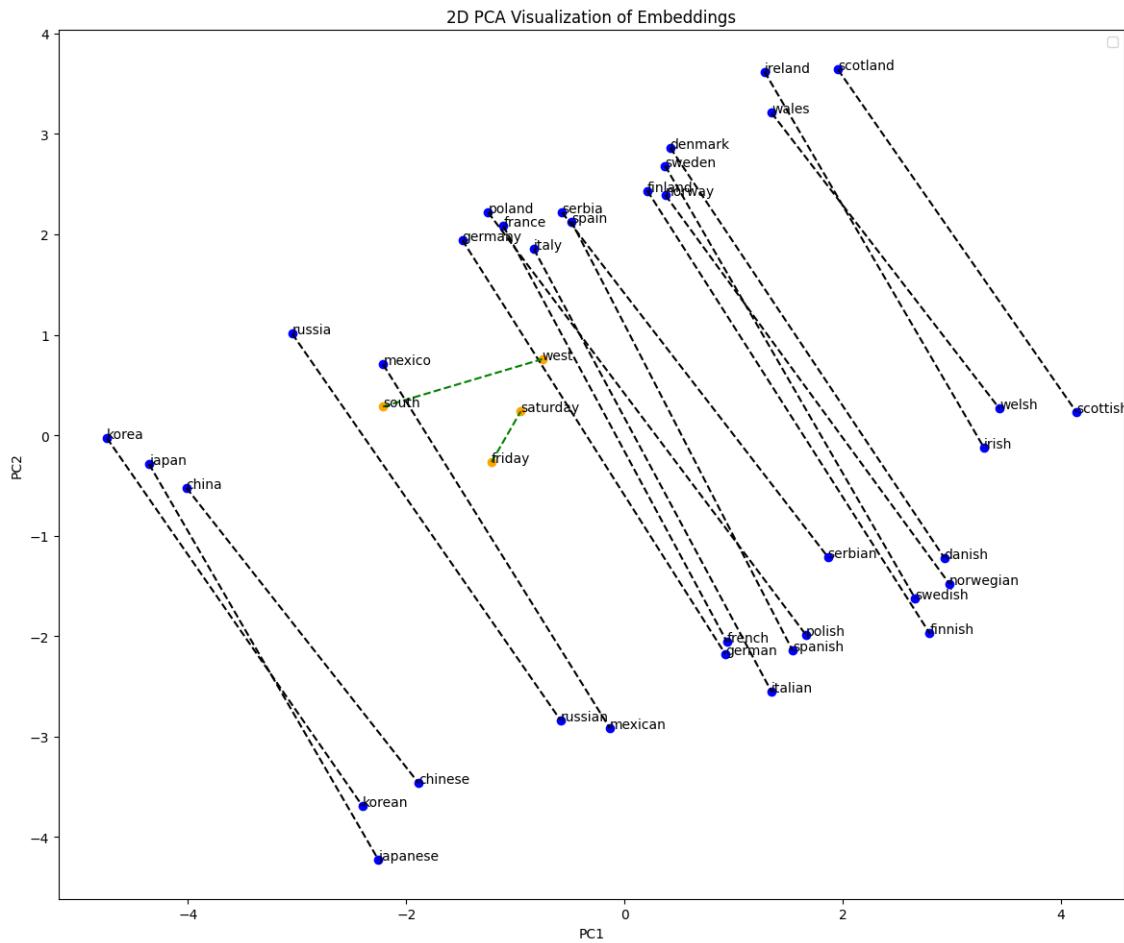
1.5
0.3
1.1
1.5
5.2
7.7
5.2
3.2
1.7
0.4
0.2
0.6
:
0.8

5.0
2.3
3.8
8.6
7.1
1.1
7.3
4.4
3.5
4.2
5.3
1.5
:
9.8



### 3D Visualization of Vector Embeddings





```

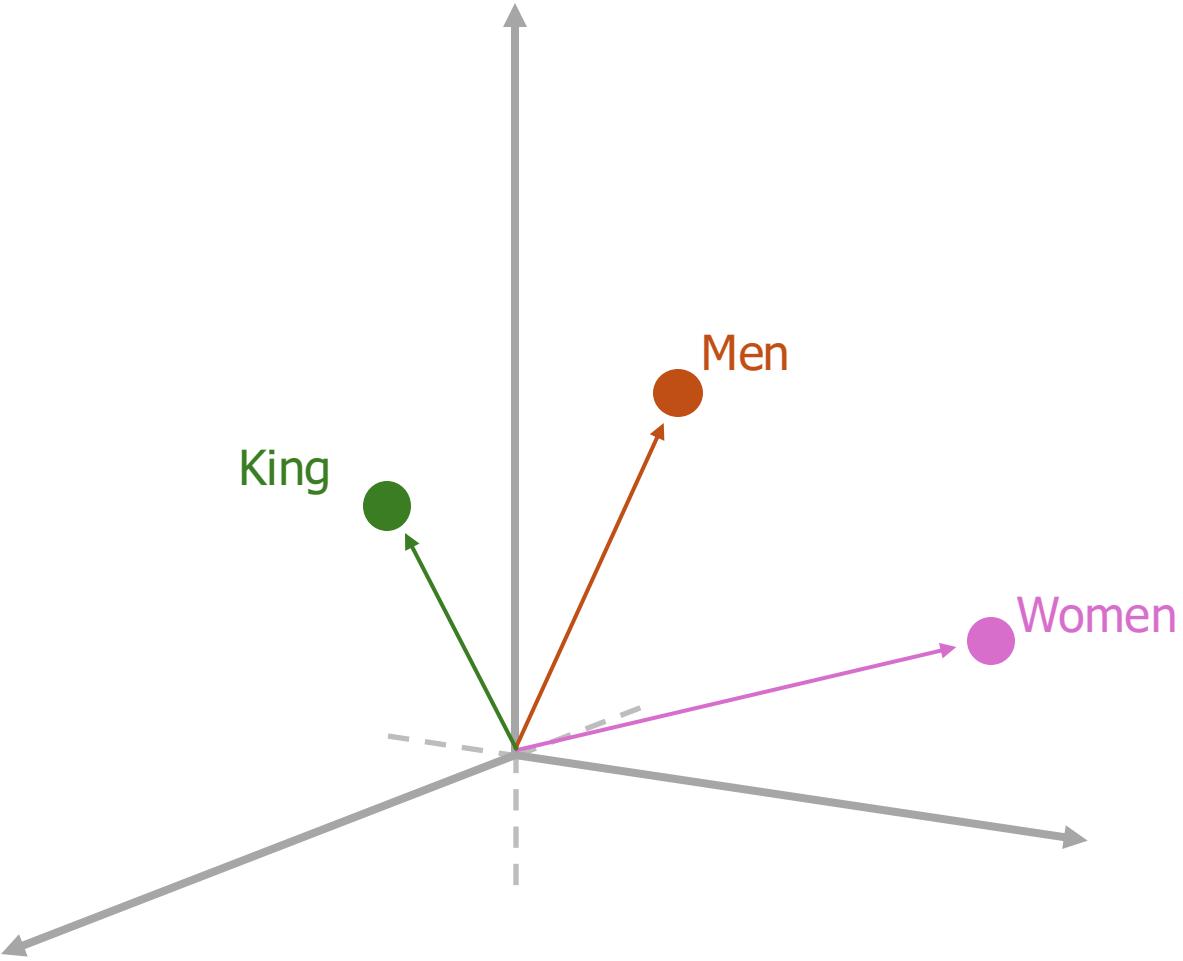
import gensim.downloader
w2v = gensim.downloader.load('glove-wiki-gigaword-300')

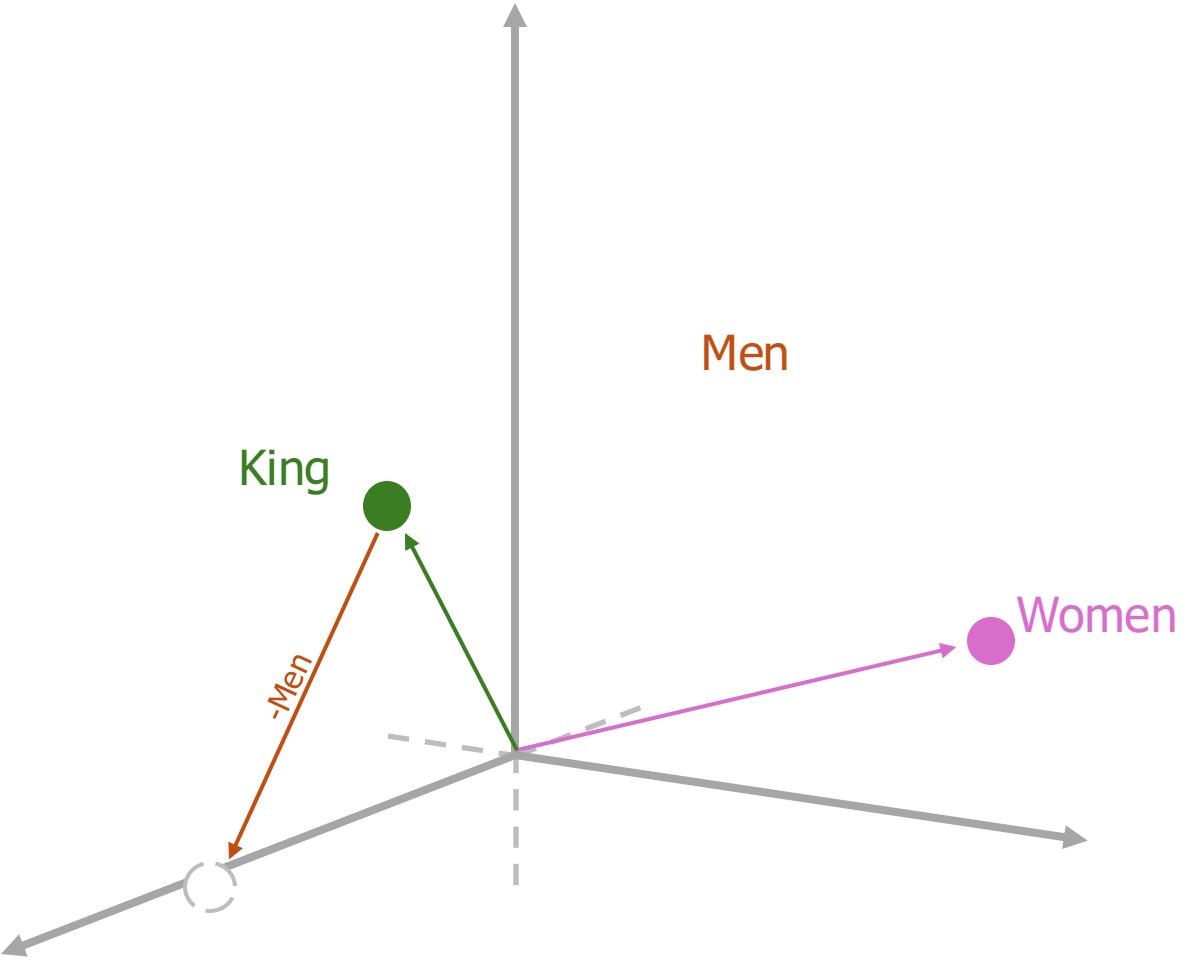
# Quantify semantic relationship between
# "germany" and "german"
relation = w2v["germany"] - w2v["german"]

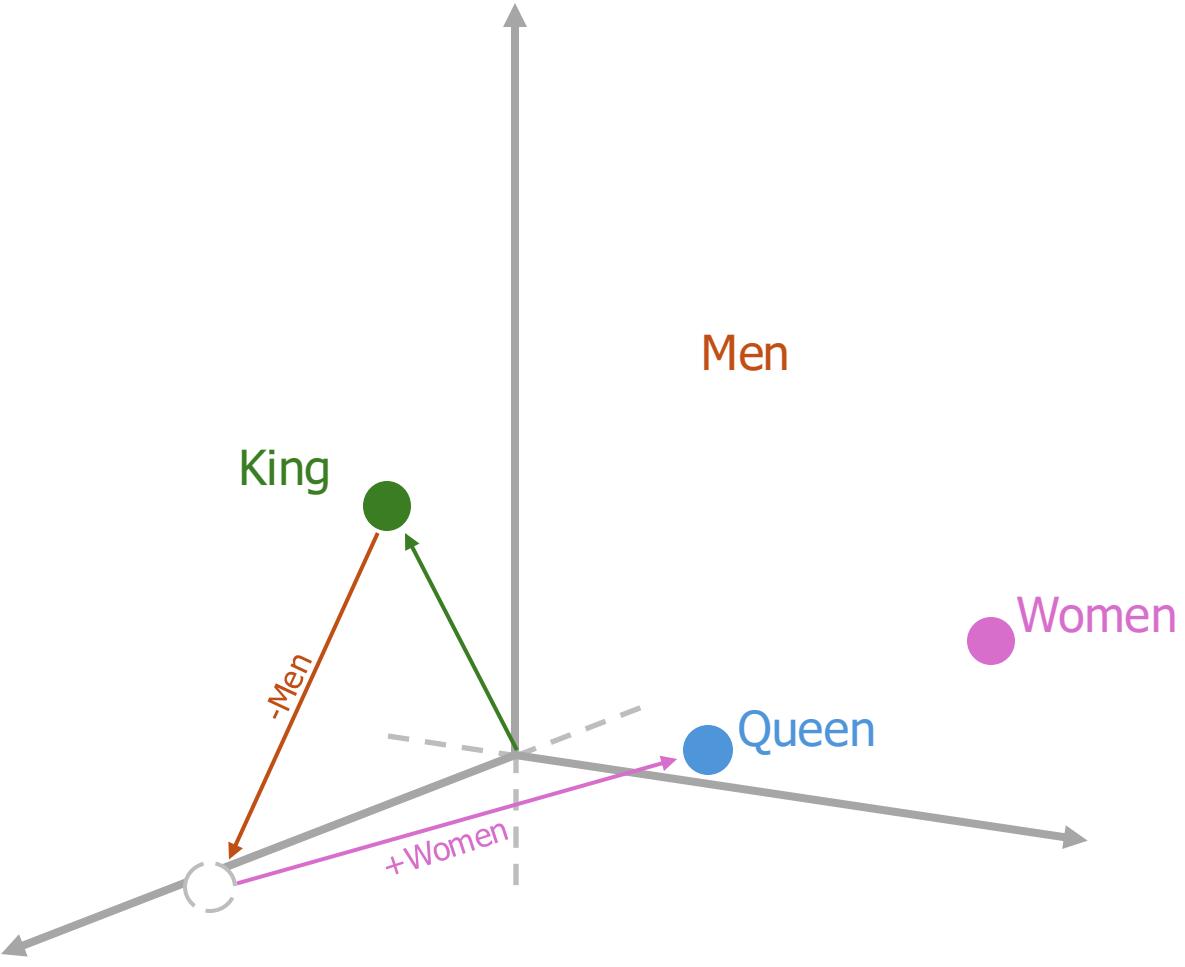
# Apply relation to different embedding
prediction = w2v["austrian"] + relation
word, similarity = w2v.most_similar(
    positive=[prediction],
    topn=1)[0]

print(f'word: "{word}" with {similarity} similarity')
# -> word: "austria" with 0.8964902758598328 similarity

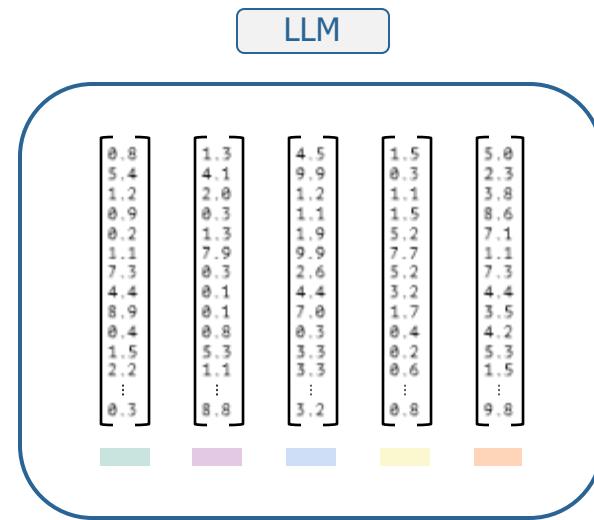
```



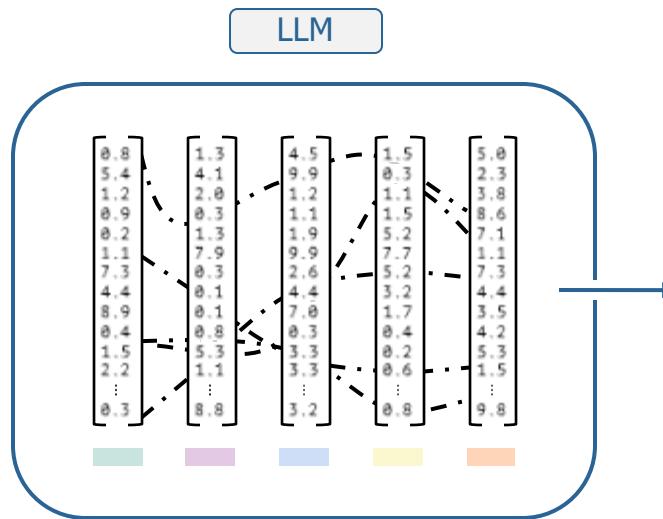




The cute green dragon trot

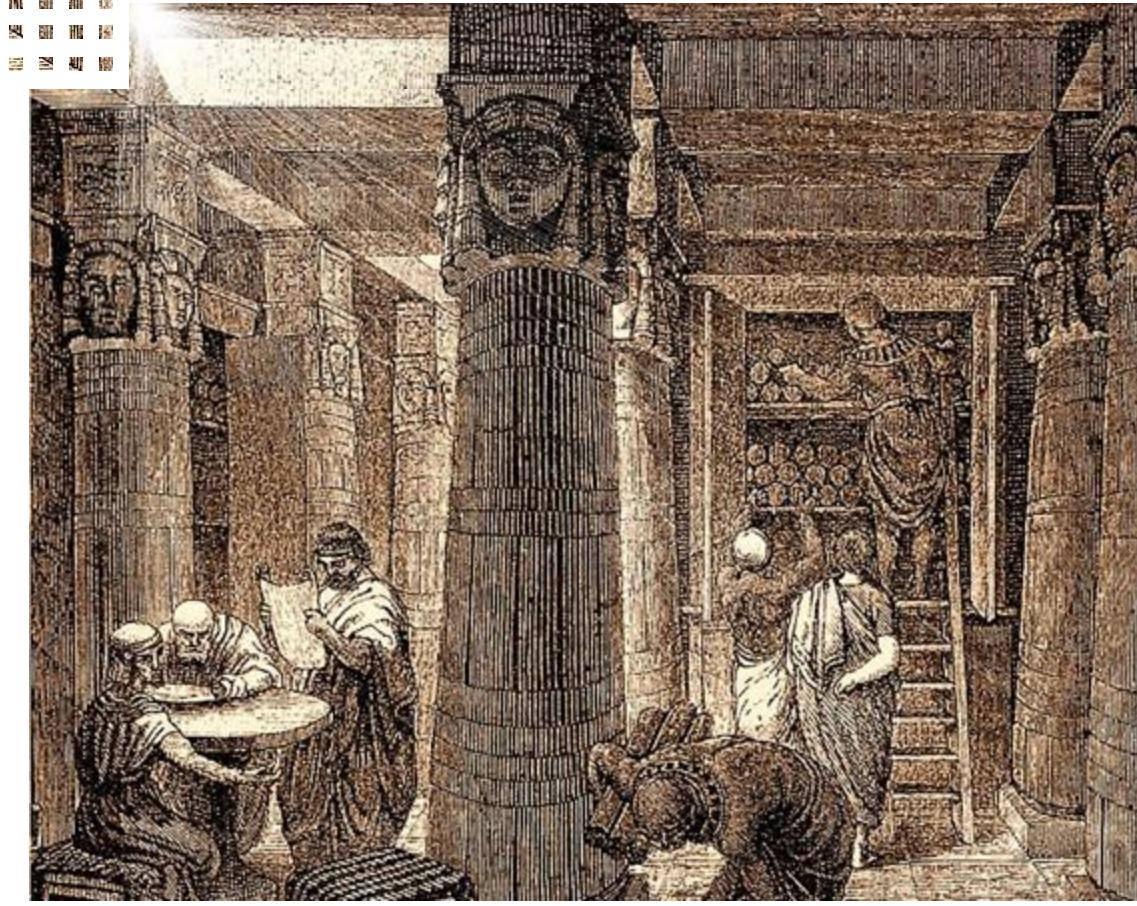


The cute green dragon trot

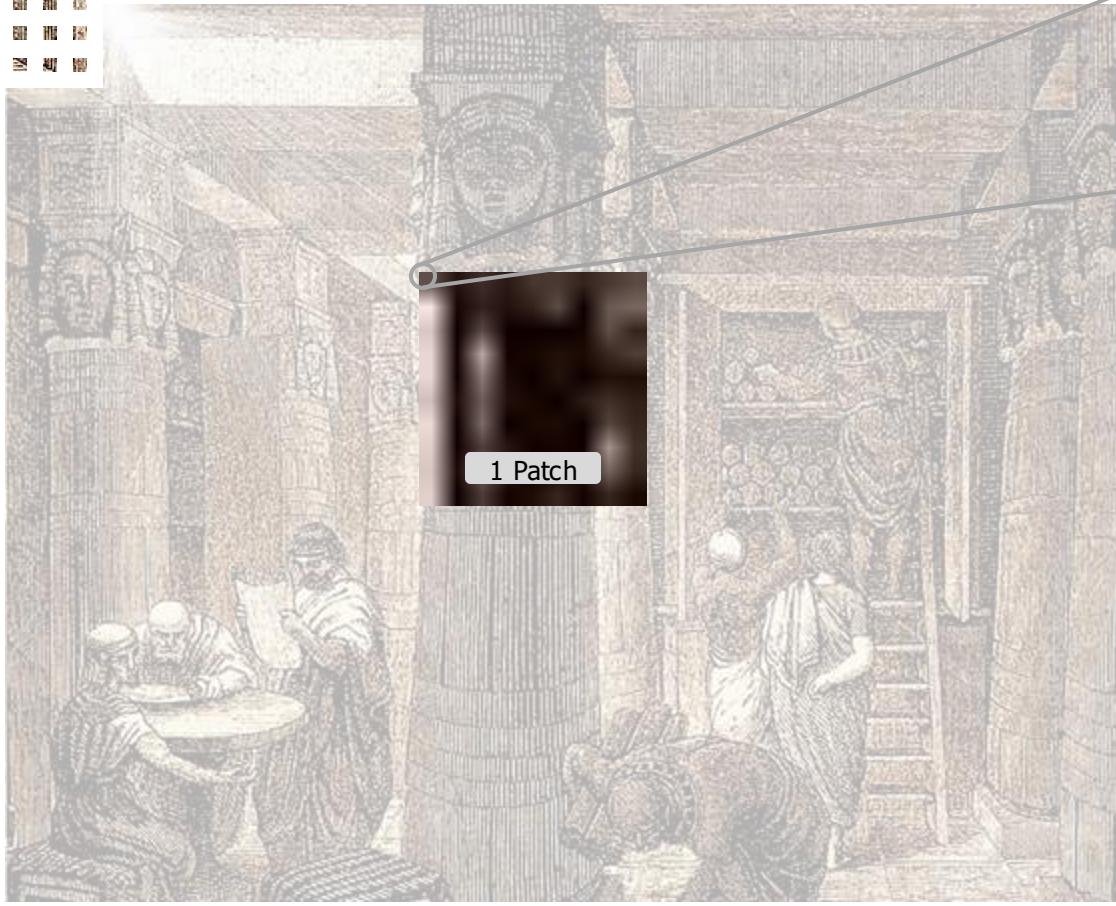


- ted      8%
- tline    5%
- tting    5%
- yl       5%
- toir    2%
- ‘       1%





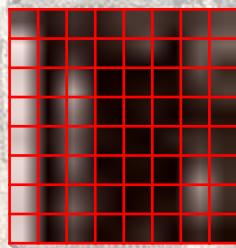
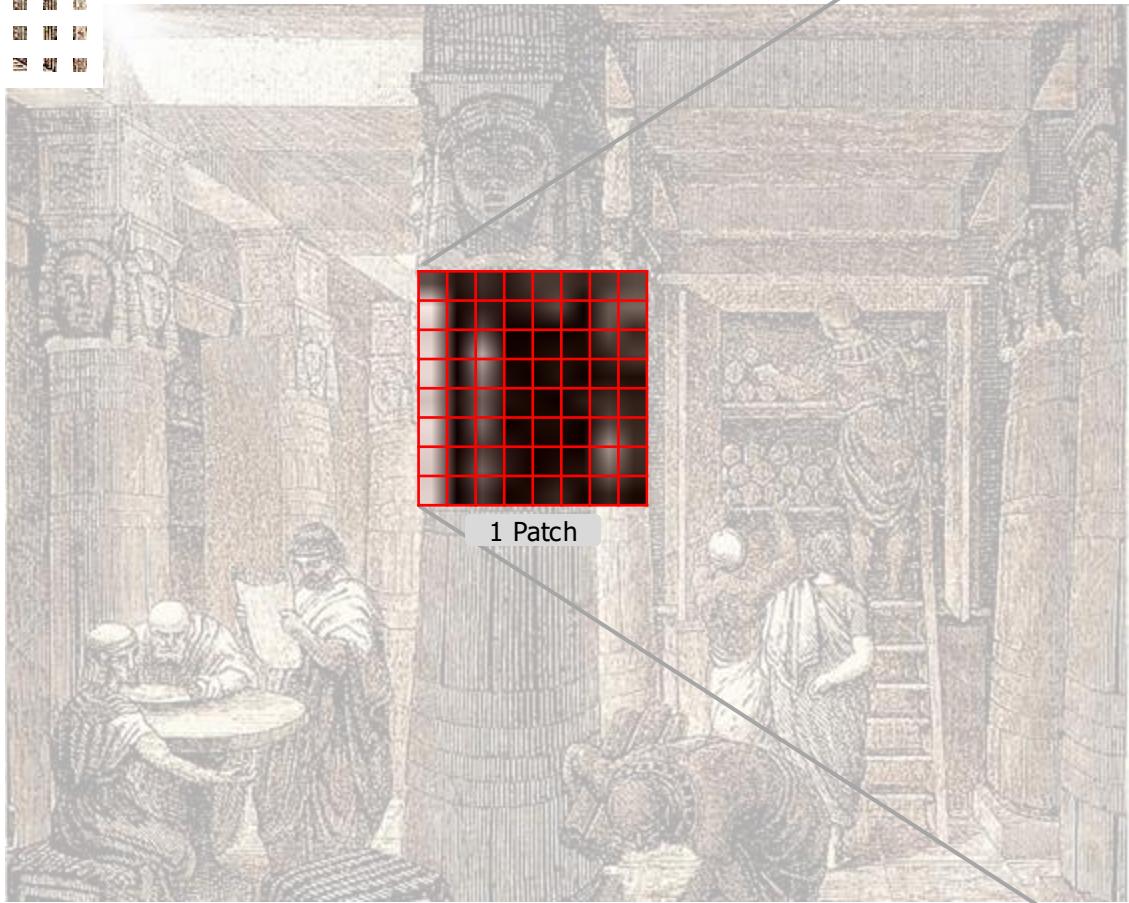




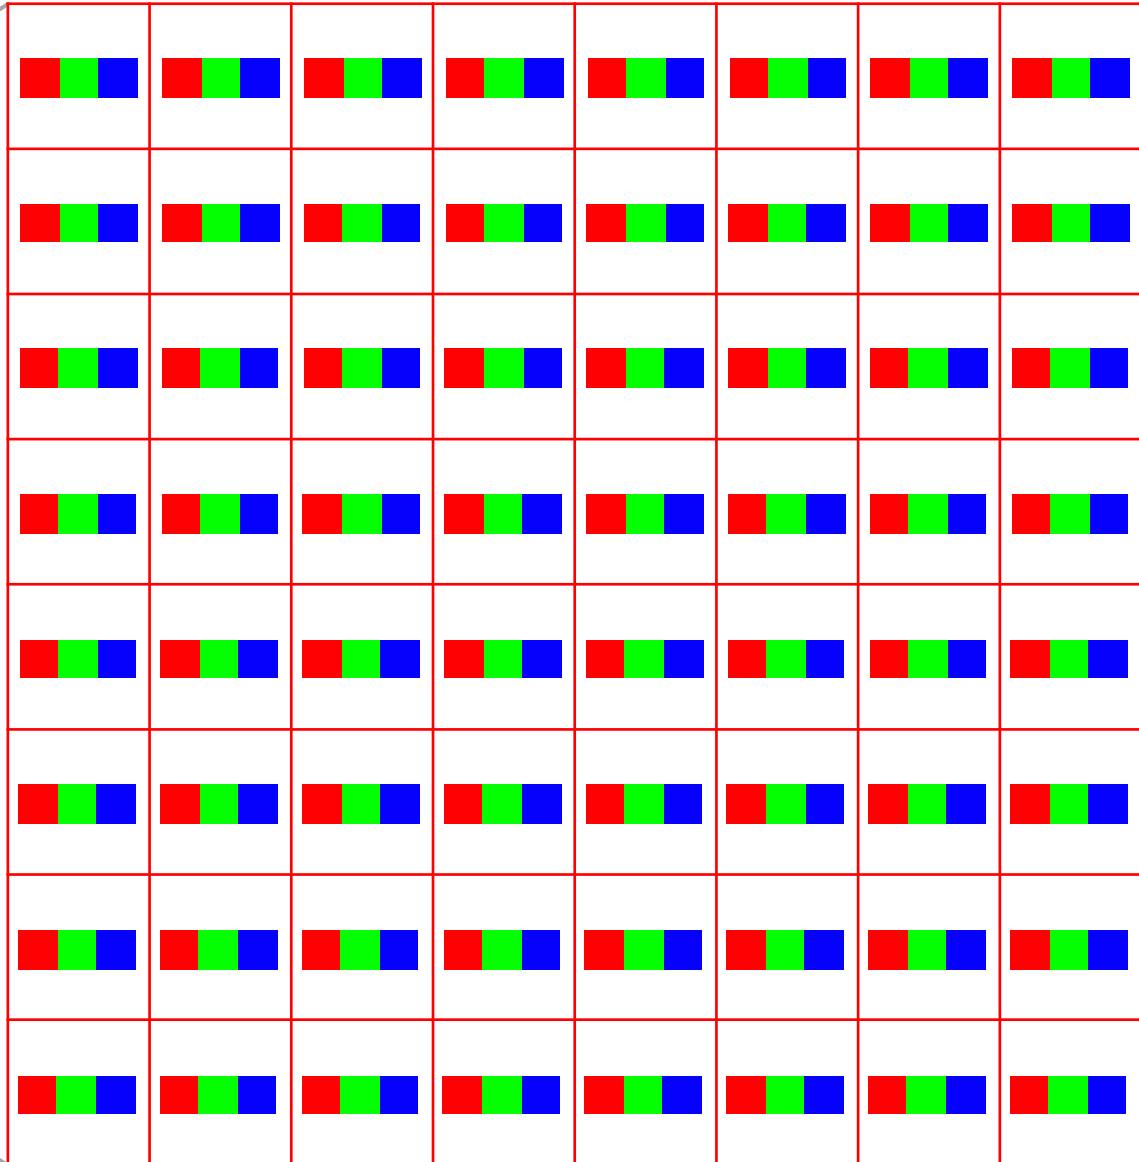
1 Patch

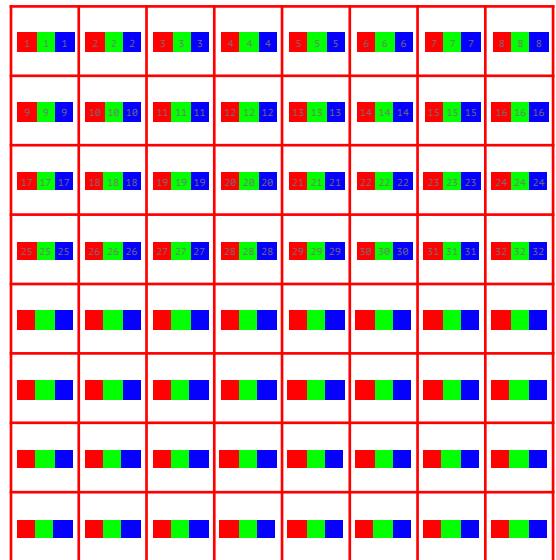
Single Pixel ( 97 88 46 )  
R G B



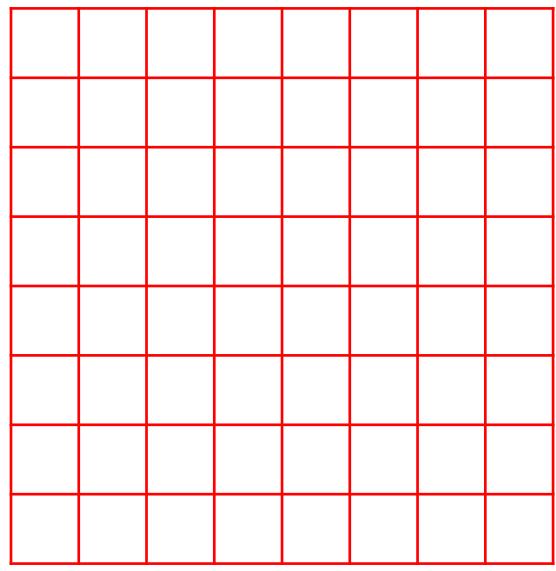


1 Patch





1 Patch



1 Patch

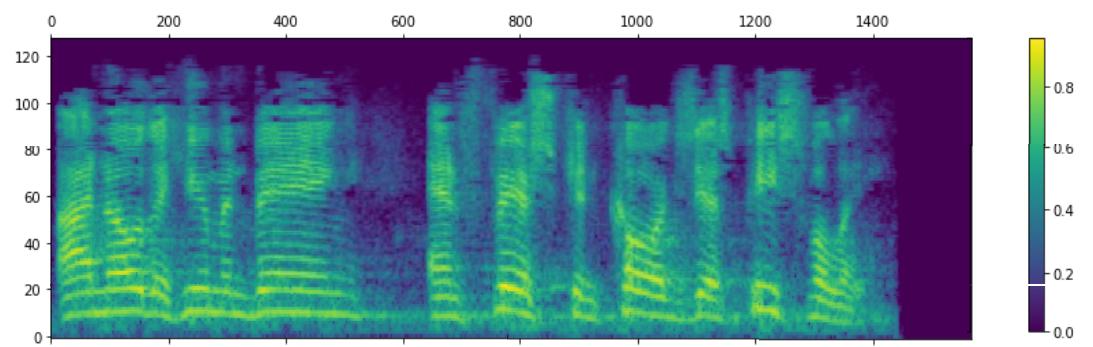


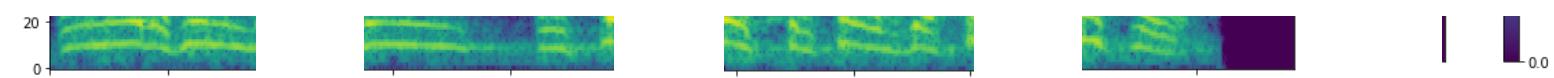
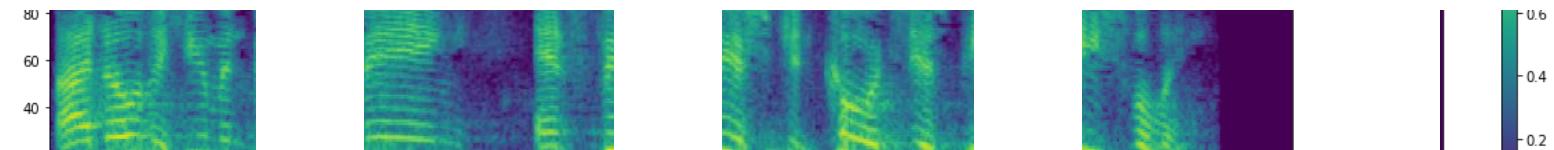
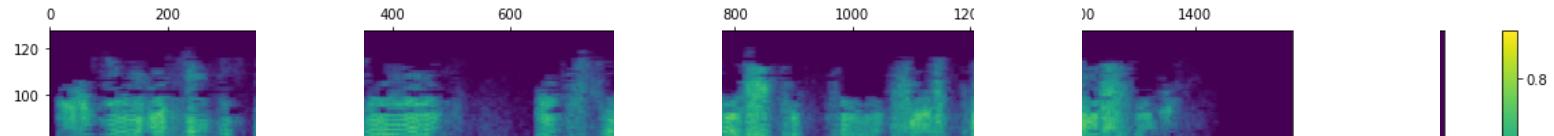
Embedding



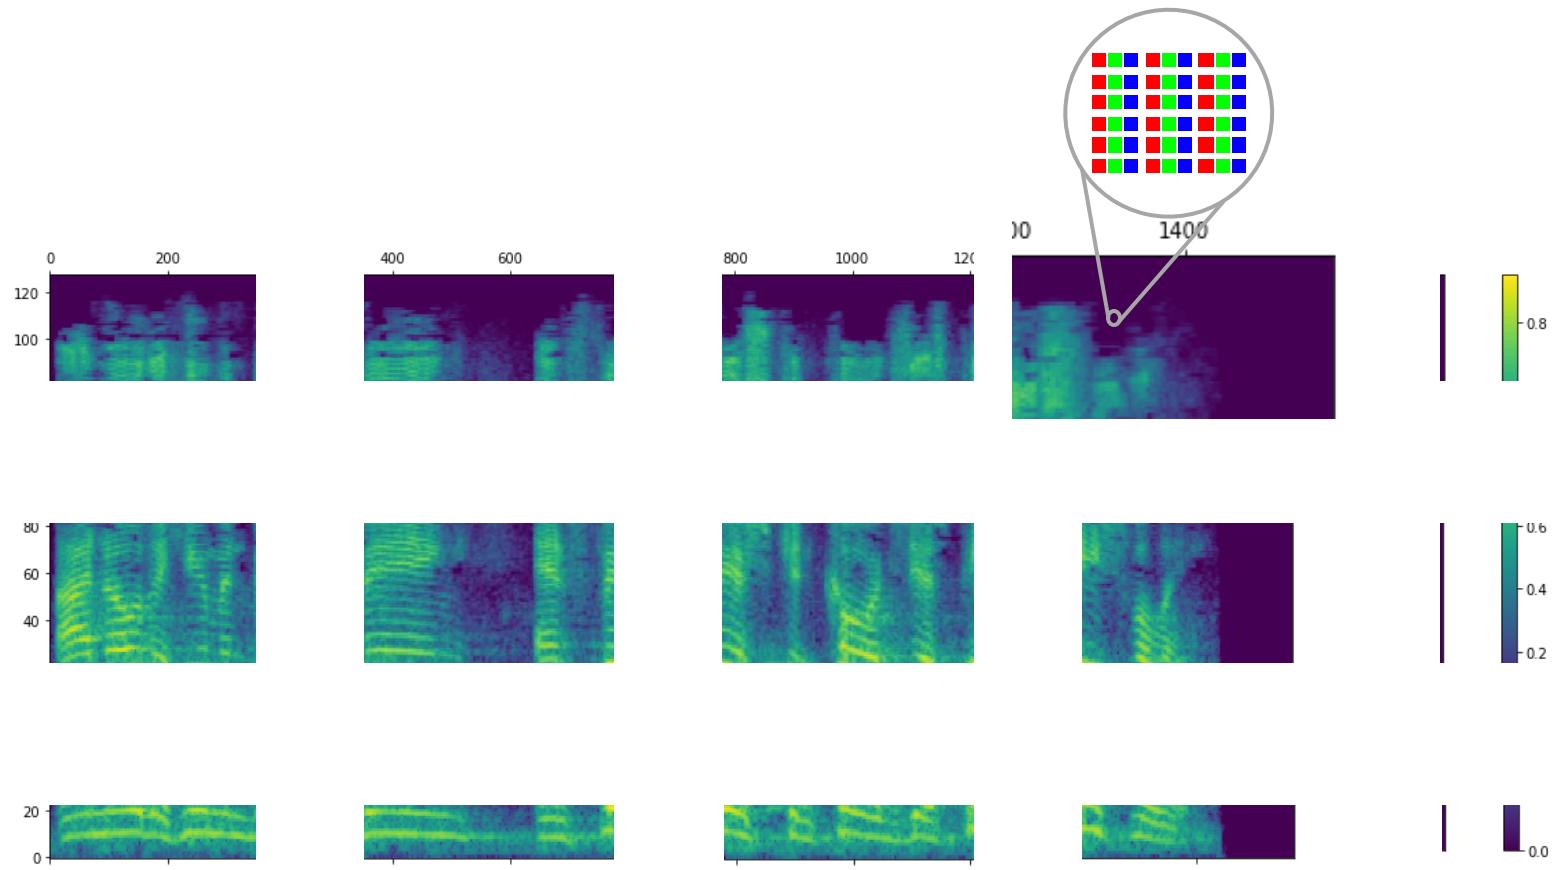


$$(\mathcal{F}f)(y) = \frac{1}{\sqrt{2\pi}^n} \int_{\mathbb{R}^n} f(x) e^{-iy \cdot x} dx$$





$$(\mathcal{F}f)(y) = \frac{1}{\sqrt{2\pi}^n} \int_{\mathbb{R}^n} f(x) e^{-iy \cdot x} dx$$

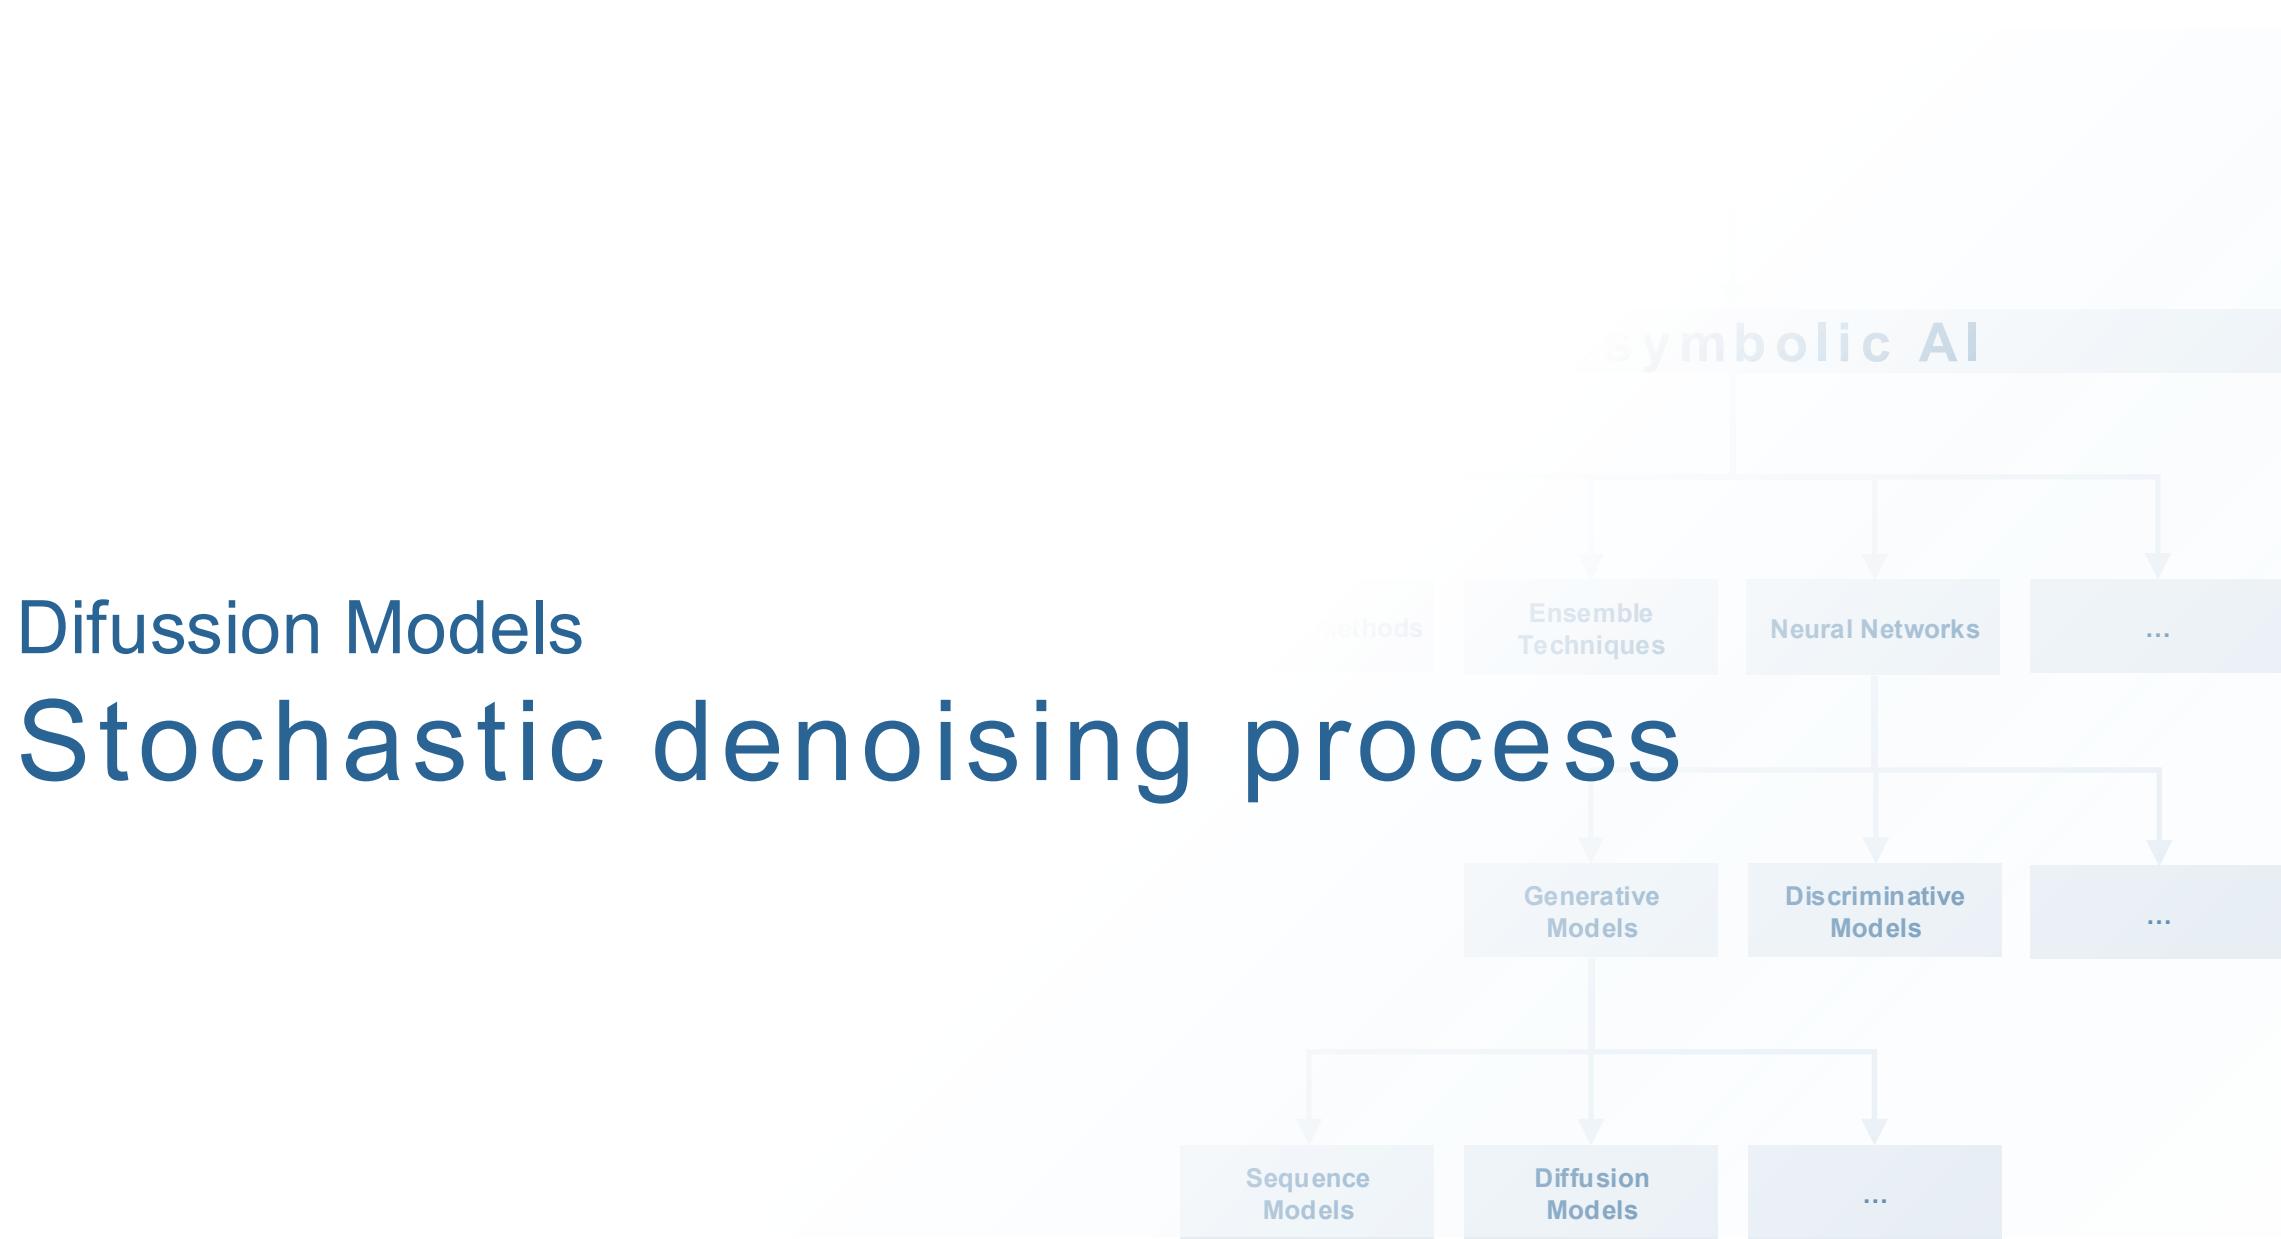


$$(\mathcal{F}f)(y) = \frac{1}{\sqrt{2\pi}^n} \int_{\mathbb{R}^n} f(x) e^{-iy \cdot x} dx$$

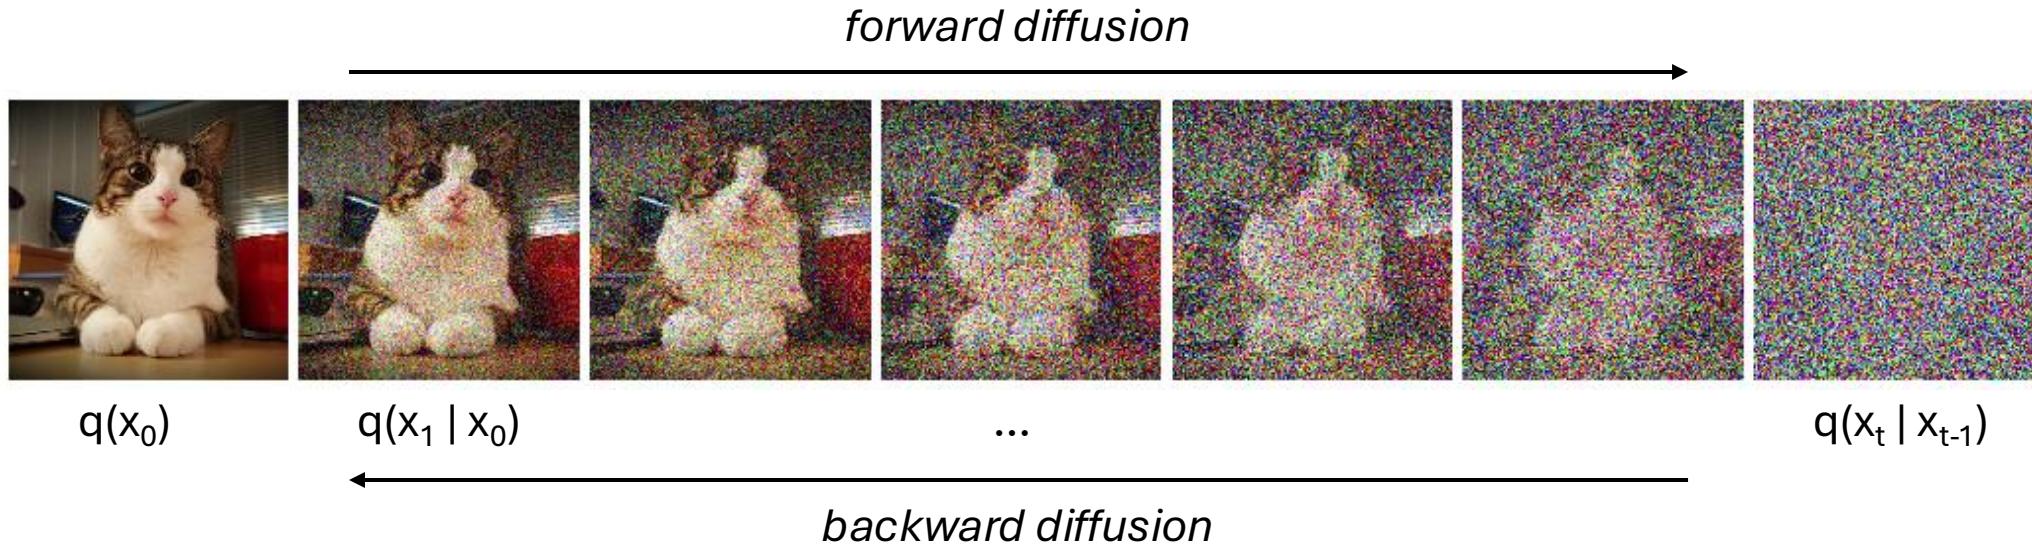
# Generative Sequence Modeling Recap

Difussion Models

# Stochastic denoising process



 Diffusion in the context of diffusion models is a **gradual transformation** that **adds random noise** (e.g. Gaussian noise) **to data** (like images) eventually converting the original data into pure noise.



$$q(x_t | x_{t-1}) = N(x_t; \sqrt{1 - \beta_t}x_{t-1}, \beta_t I)$$

with  $\beta_t \in (0, 1)$

“Create a normal distribution that centers around  $\sqrt{1 - \beta_t}x_{t-1}$  with variance  $\beta_t I$ ”

$$Insight: q(x_t | x_0) = N(x_t; \sqrt{a_t}x_0, (1 - a_t)I)$$

with  $a_t = \prod 1 - \beta_i$

“Chaining multiple Gaussian distributions, allows us to directly sample from  $x_0$  to any timestep  $t$ .”

1.  $q(x_t | x_{t-1})$  represents the noised image at step  $t$  conditioned by its previous step3. Nice mathematical introduction to diffusion models:

2. Image src: <https://cvpr2022-tutorial-diffusion-models.github.io/>

<https://lilianweng.github.io/posts/2021-07-11-diffusion-models/>

# Introduction To Generative AI Recap

## Symbolic vs. Subsymbolic AI

- **Symbolic AI** defines a set of formal, humanly understandable symbols and employs explicit rules to draw logical conclusions.
- **Subsymbolic AI** leverages mathematical models to learn patterns from data without requiring hand-crafted rules.

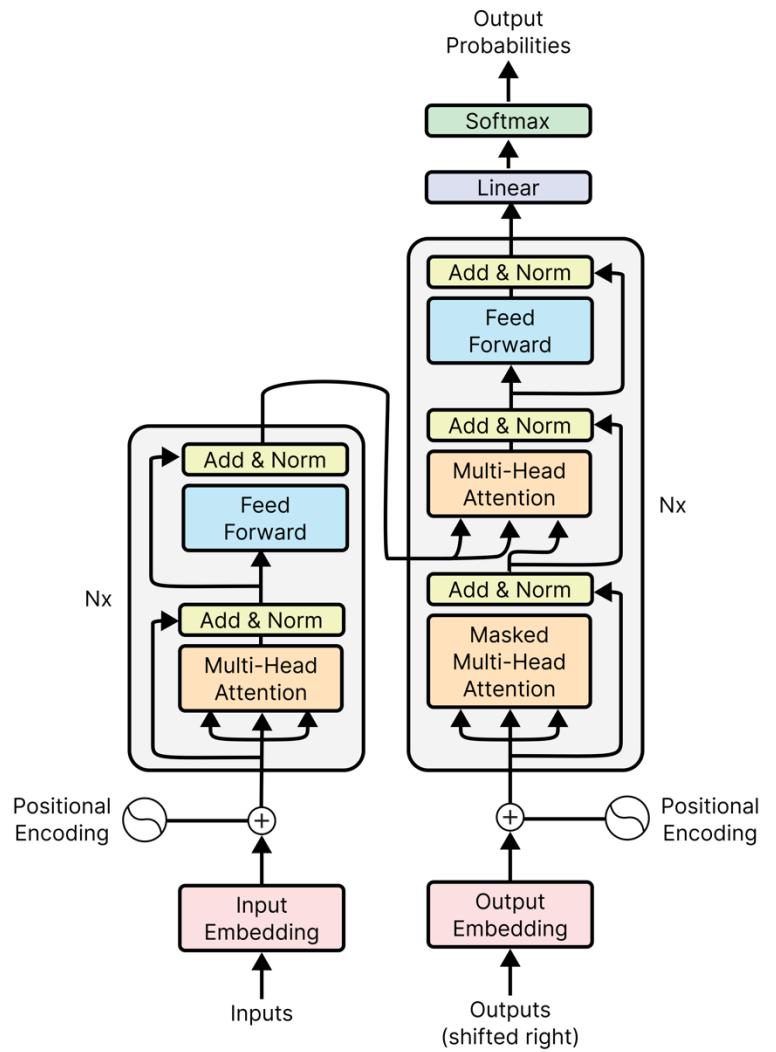
## Discriminative vs. Generative Models

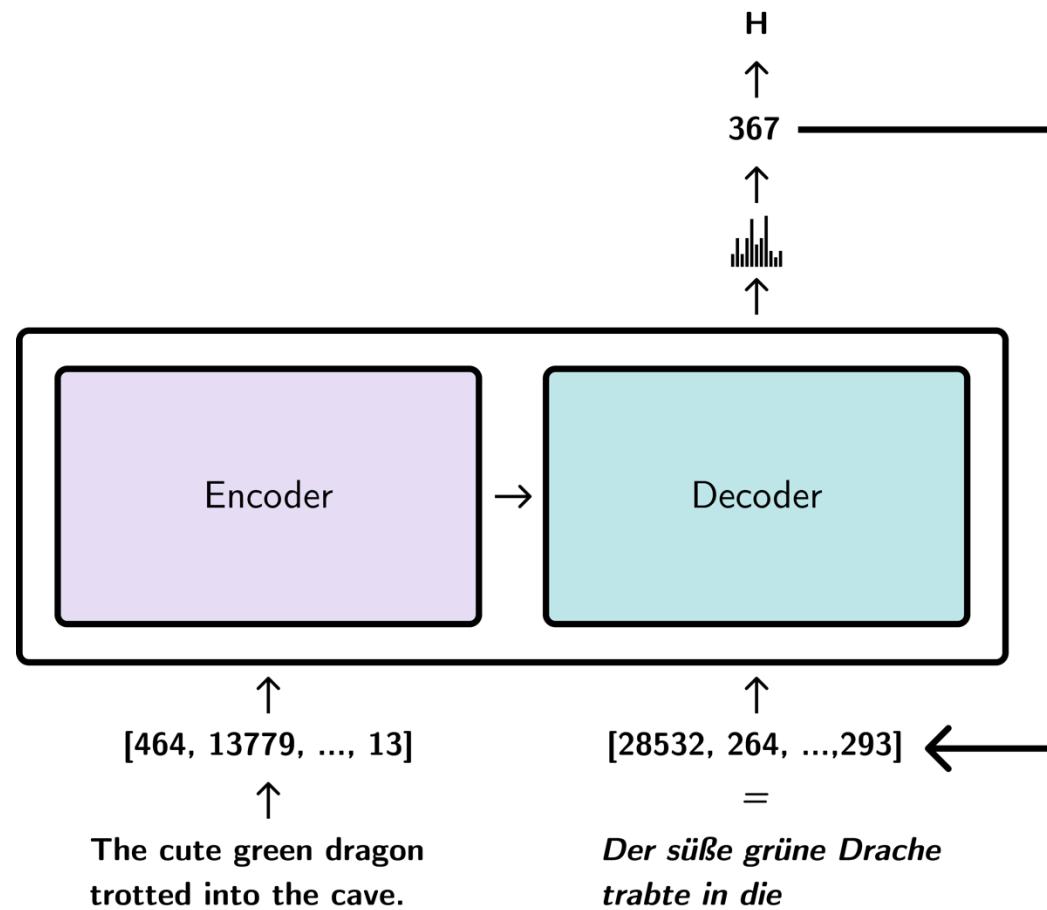
- **Discriminative Models** learn the conditional probability  $p(Y|X)$ , modeling the decision boundary between classes or the mapping from inputs to outputs.
- **Generative Models** learn the joint probability  $p(X,Y)$  or, in the unsupervised case, the data distribution  $p(X)$ .

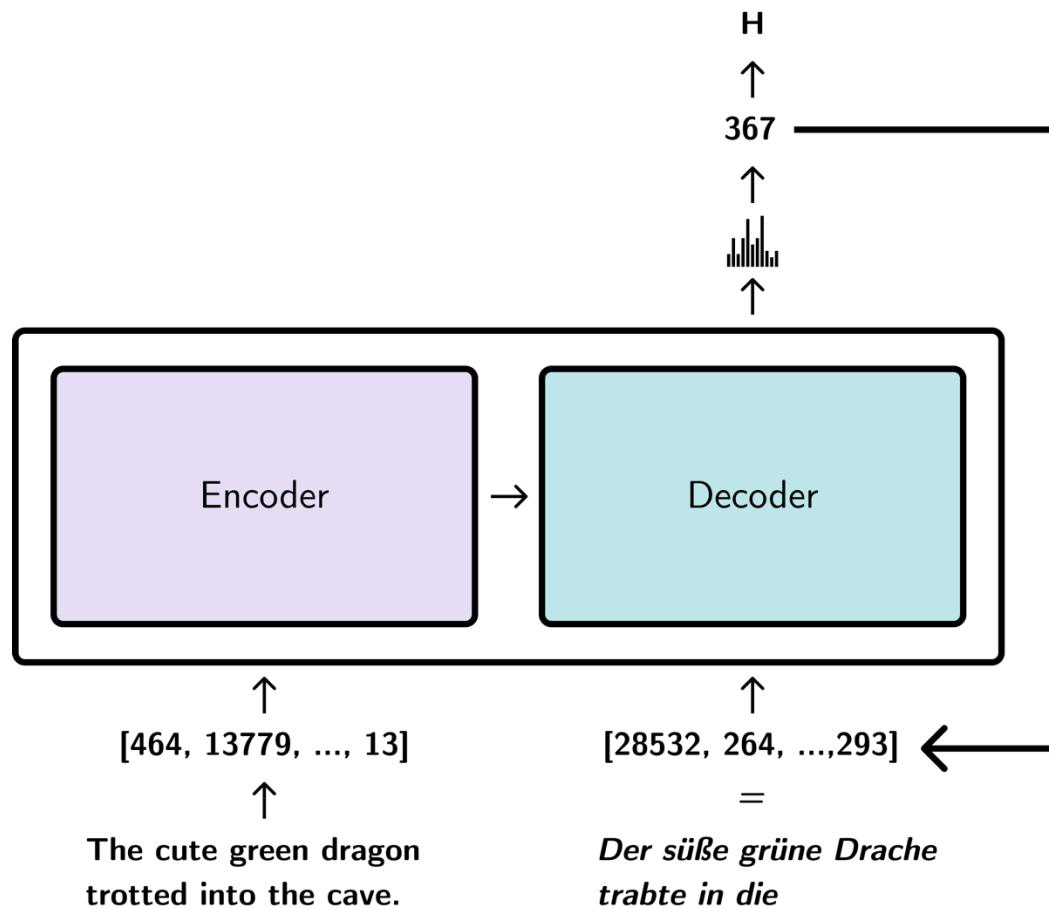
## Sequence Modeling vs. Diffusion Modeling

- **Sequence Models** process data as ordered sequences of discrete units (tokens in language models or patches in vision models). Those data atoms are mapped to continuous embeddings – dense vector representations capturing semantic or structural information.
- **Diffusion Models** learn the data distribution  $p(X)$  by training a model to iteratively denoise data, reversing a gradual noising process applied during training.

# A bird's eye view of the Transformer Architecture.

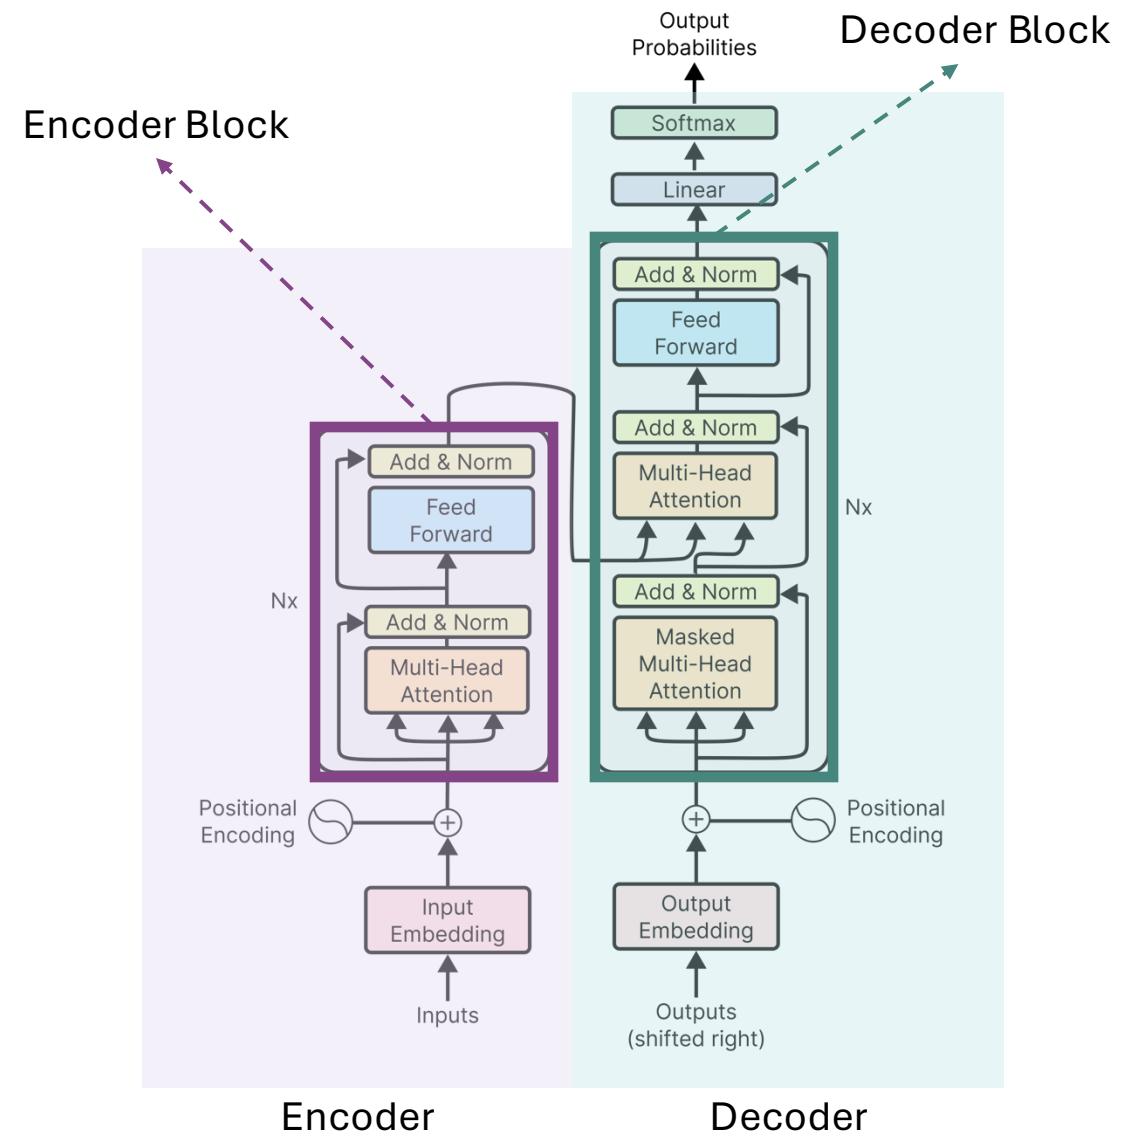
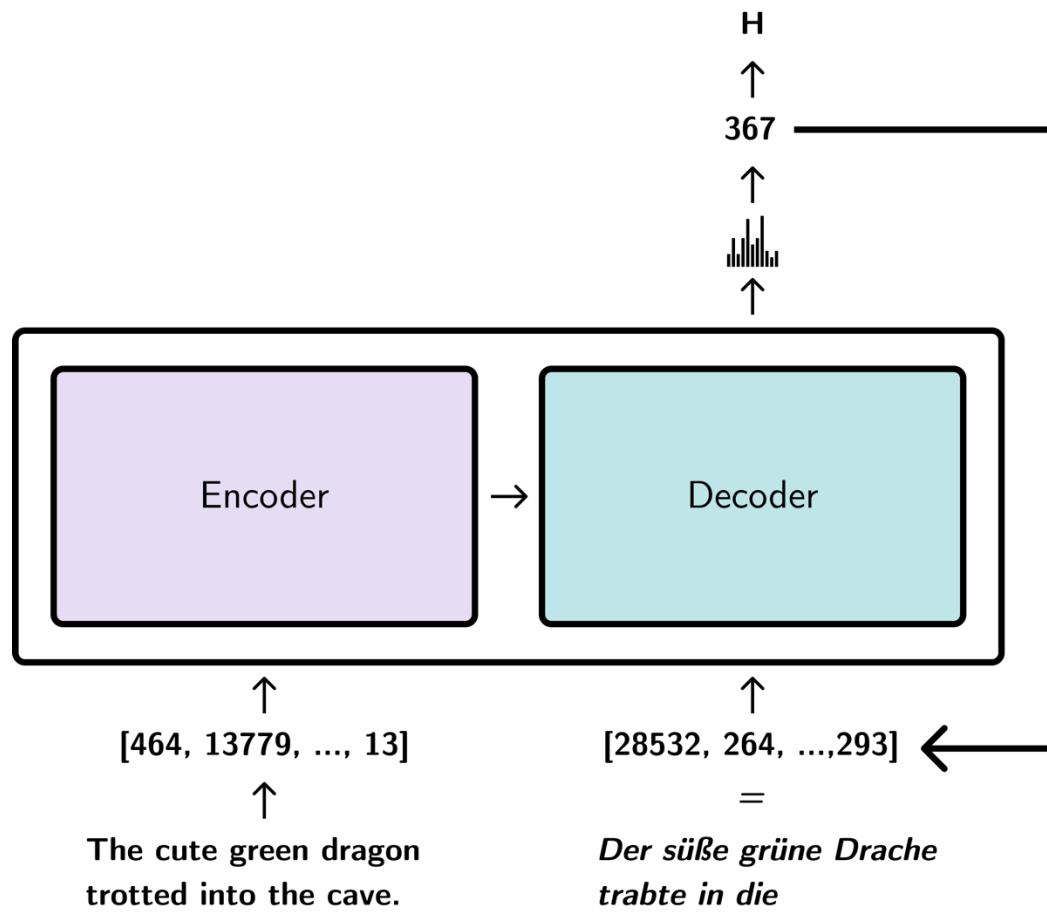


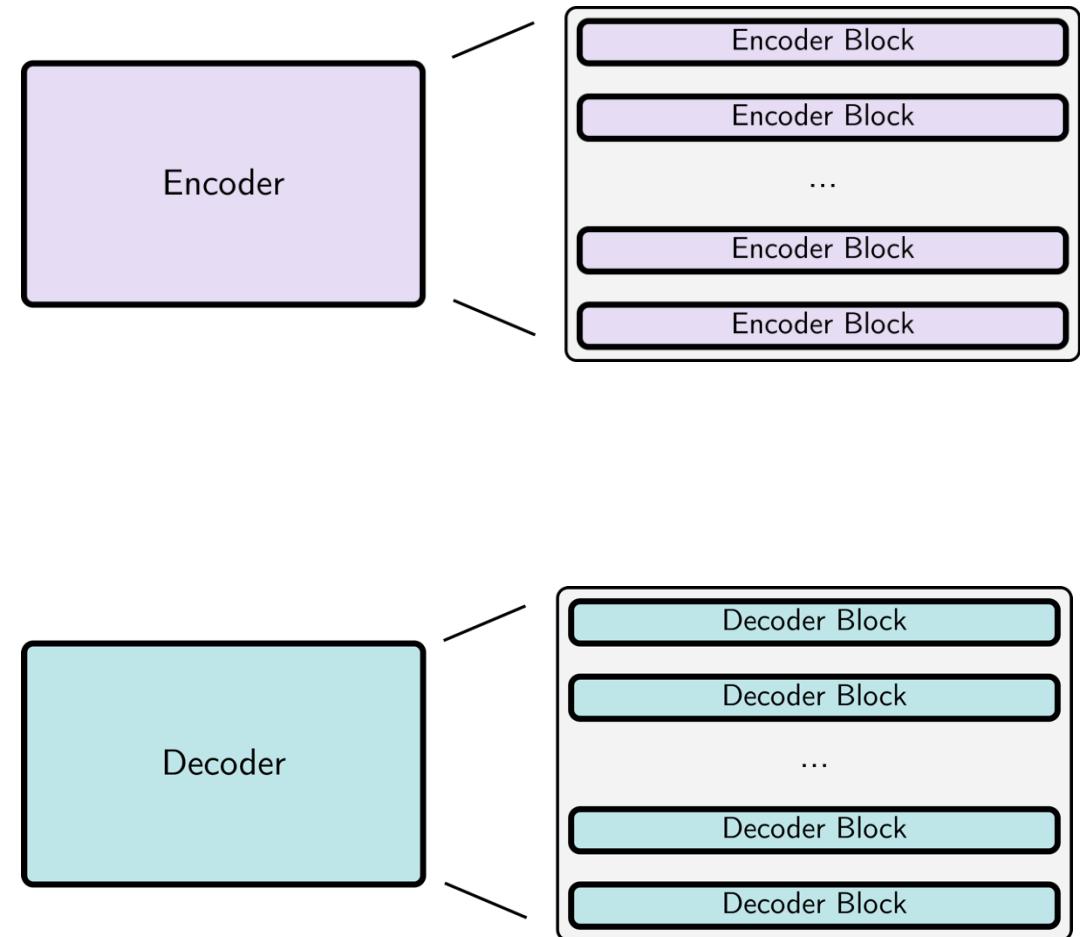
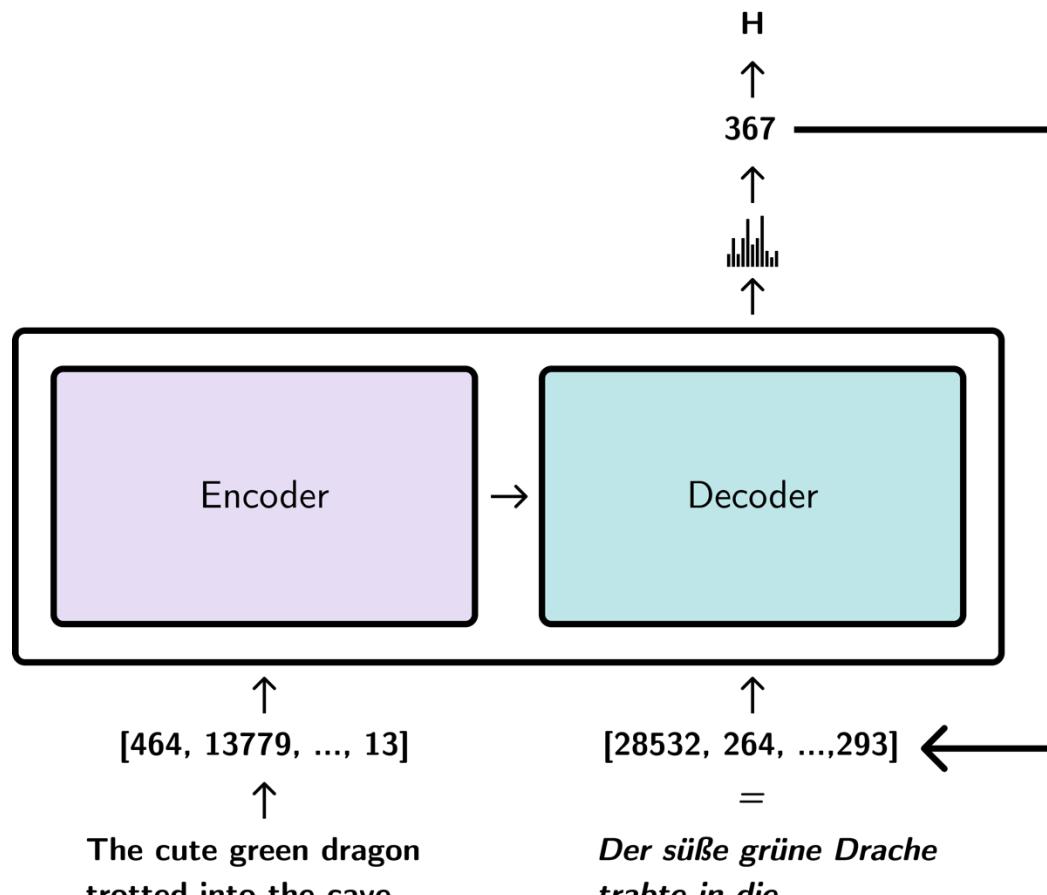


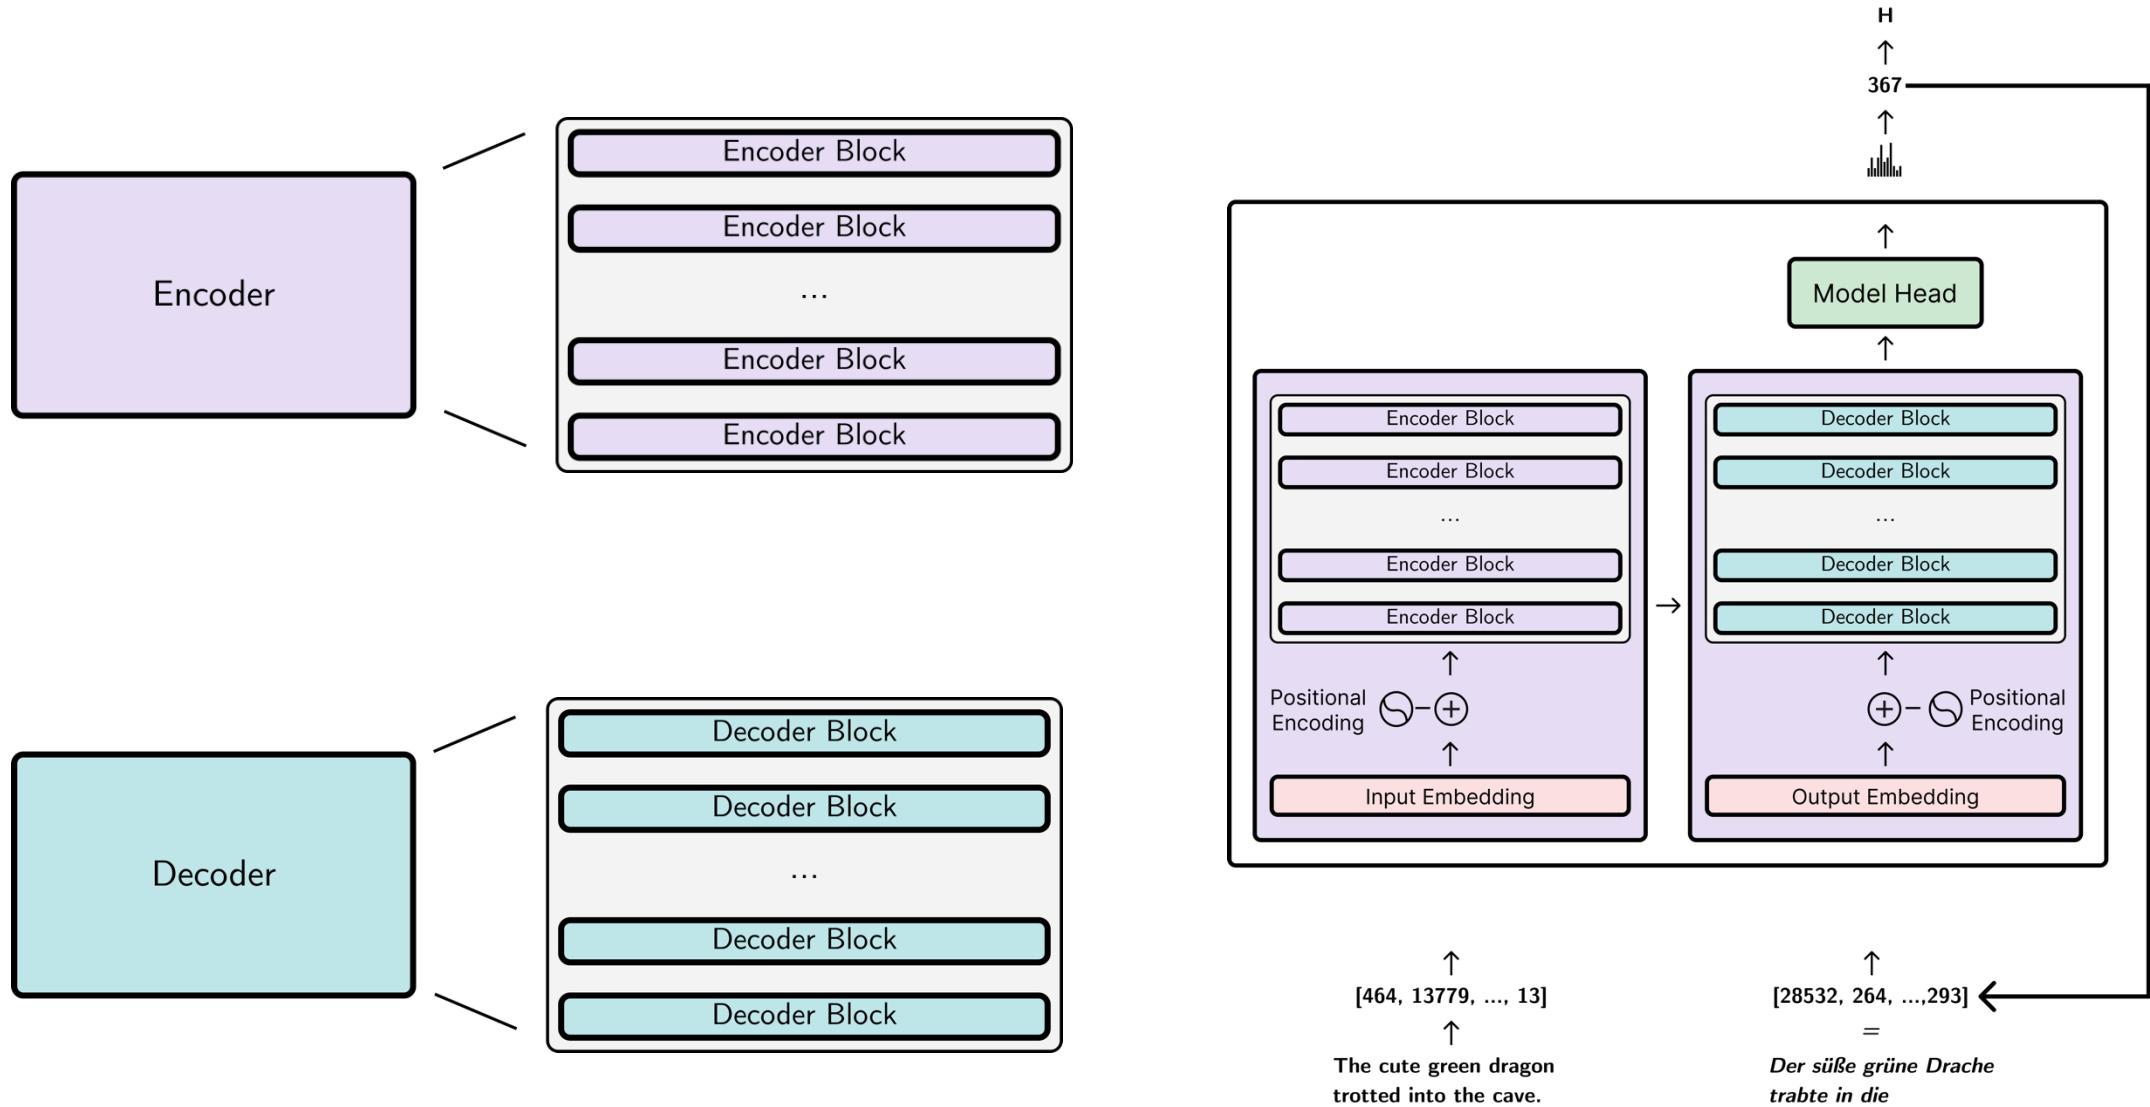


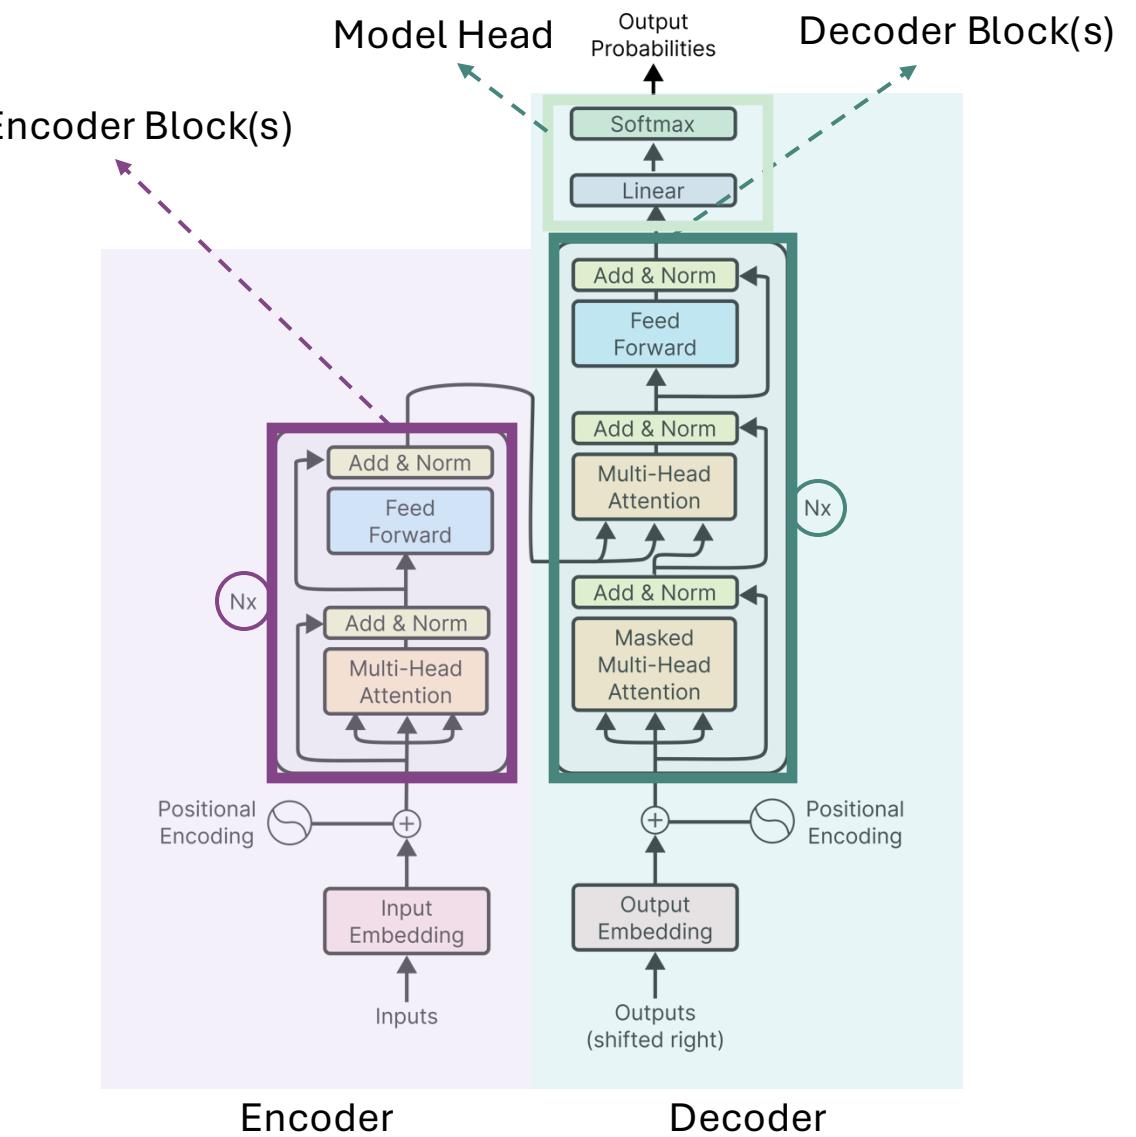
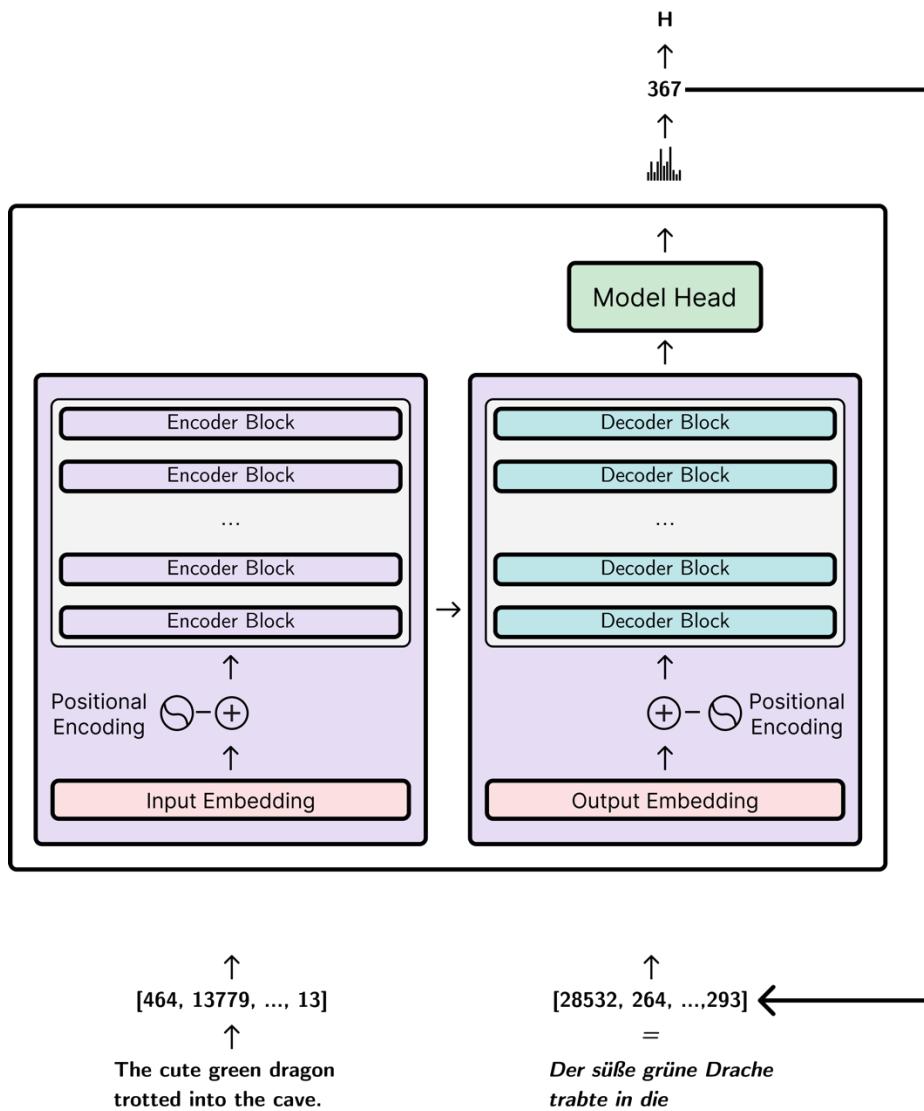
The **encoder** transforms the input sequence into a rich, contextual vector representation that captures the meaning of and relationships between elements.

The **decoder** takes the encoded input from the encoder and autoregressively generates an output sequence token by token, using previously generated tokens and the context of the encoder to produce a transformed sequence (e.g. a translation or summary).









# Positional Encoding

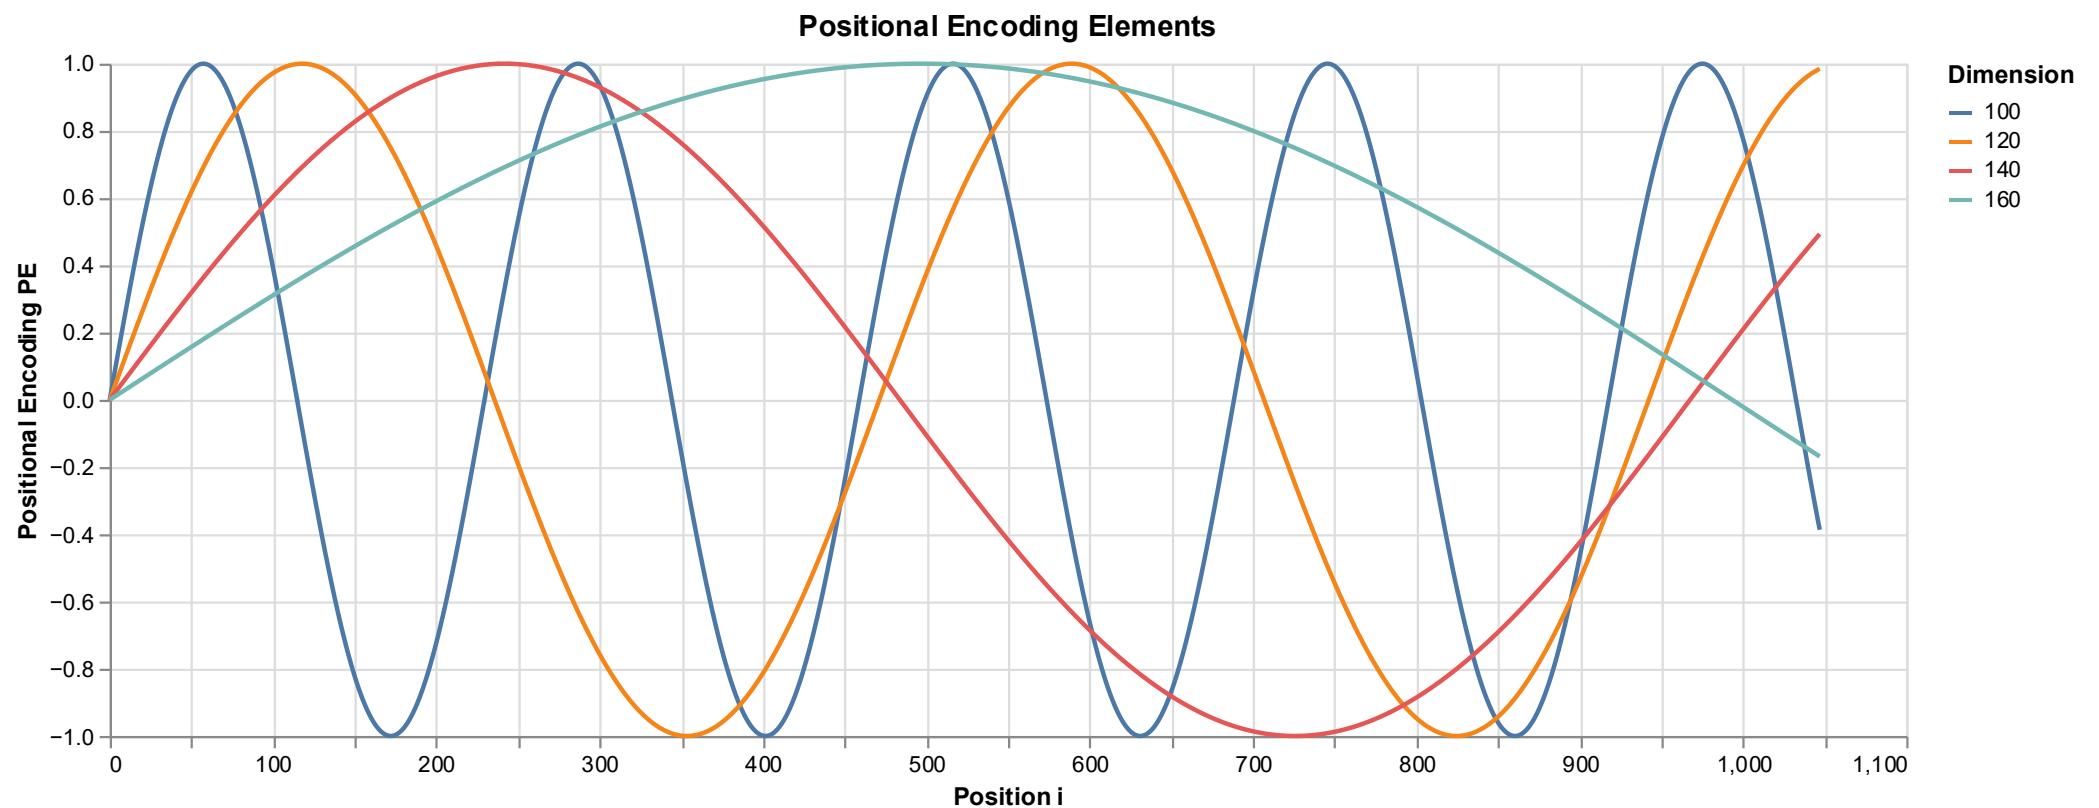
 The embedding layer and subsequent layers are inherently permutation invariant, but the order plays a critical role in sequence processing tasks.

- $PE(pos, 2_i) = \sin(pos/10000^{2i/d_{model}})$
- $PE(pos, 2_{i+1}) = \cos(pos/10000^{2i/d_{model}})$

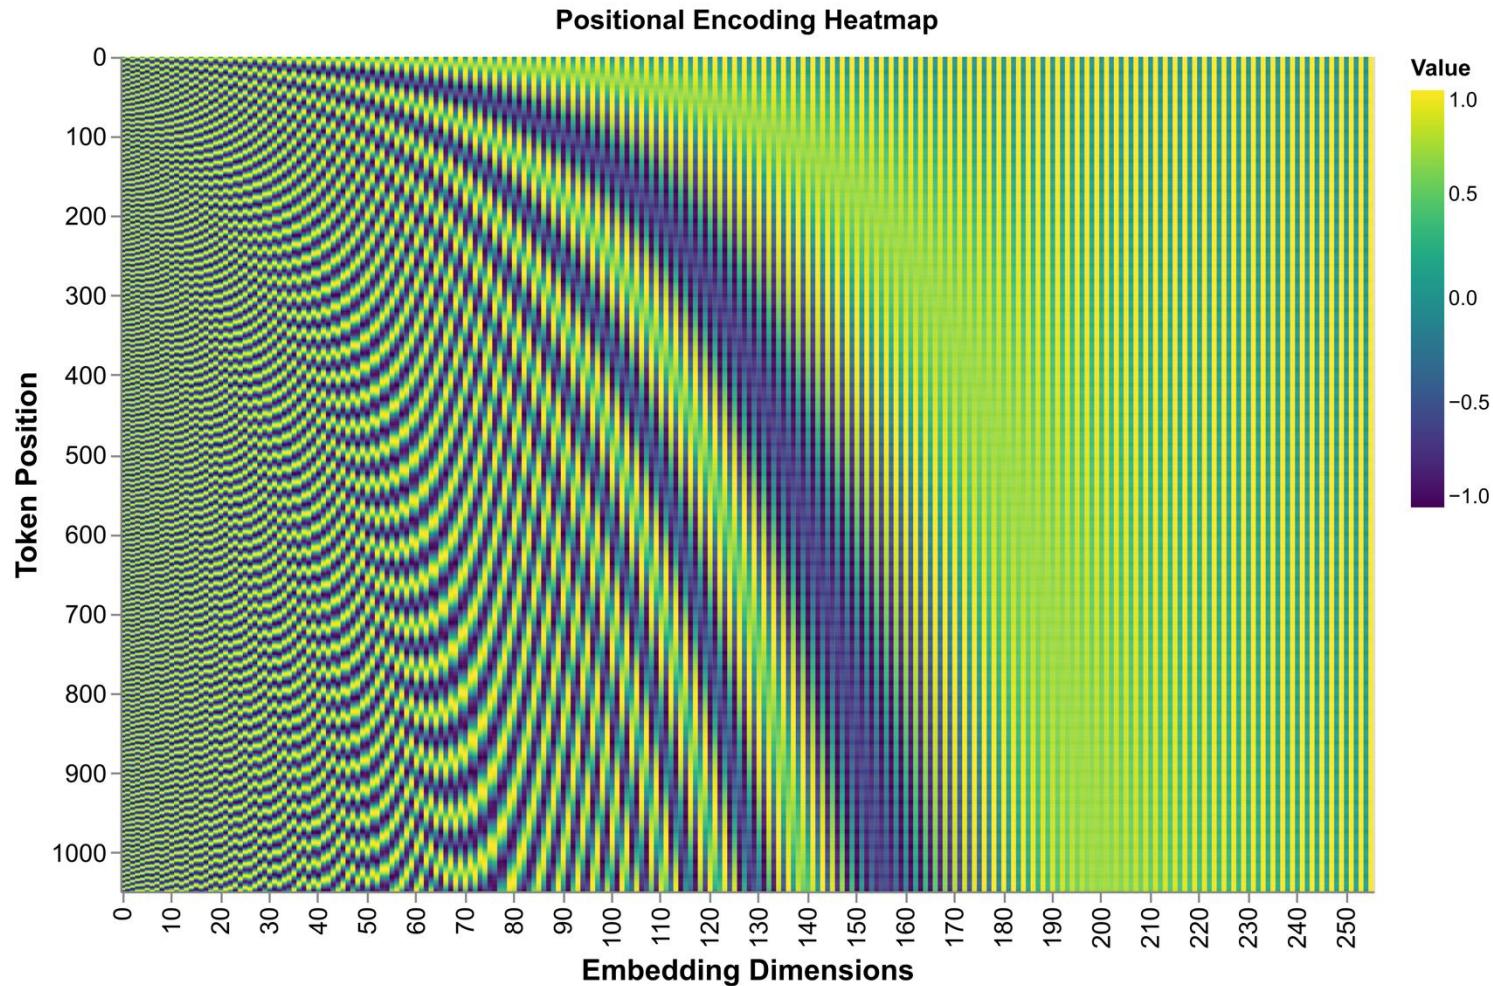


$$\overrightarrow{PE} = \begin{bmatrix} \sin(pos/10000^{0/d_{model}}) \\ \cos(pos/10000^{0/d_{model}}) \\ \sin(pos/10000^{2/d_{model}}) \\ \cos(pos/10000^{2/d_{model}}) \\ \vdots \\ \vdots \\ \sin(pos/10000^{d_{model}-2/d_{model}}) \\ \cos(pos/10000^{d_{model}-2/d_{model}}) \end{bmatrix}$$

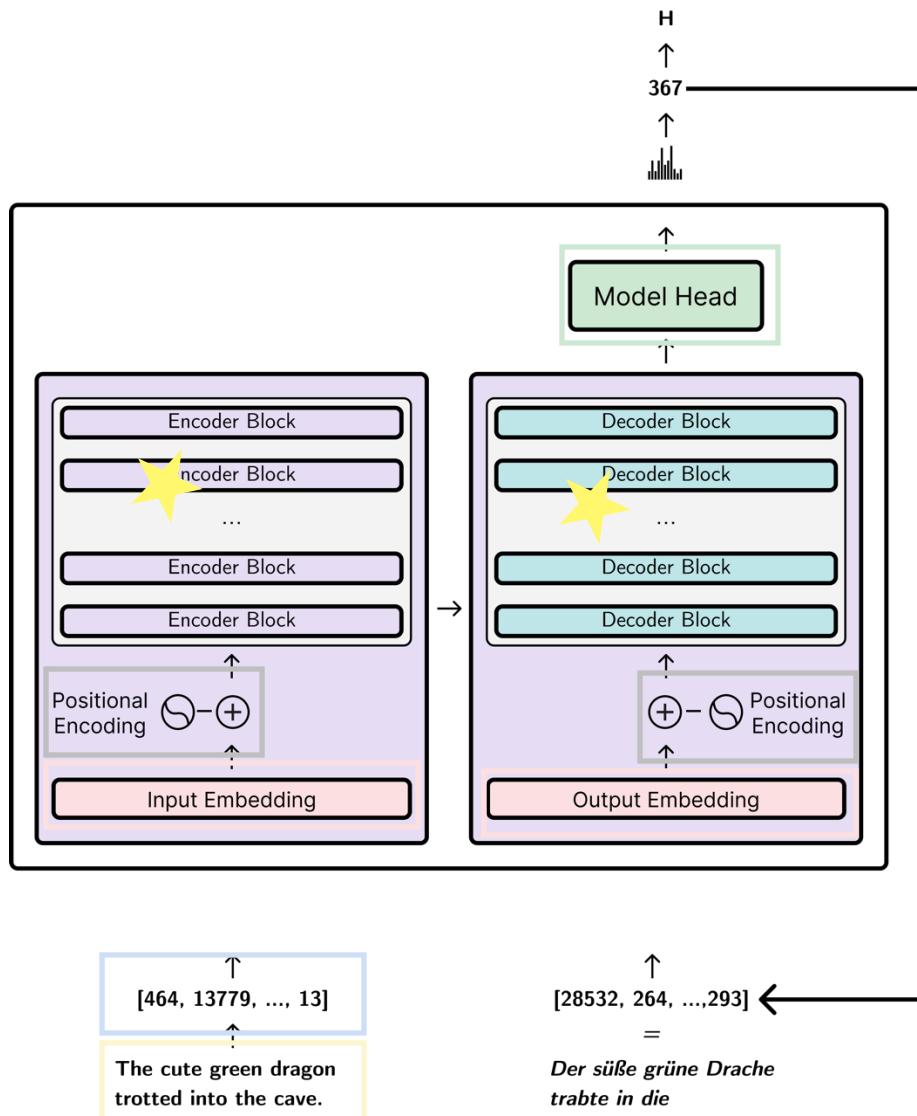
- $PE(pos, 2_i) = \sin(pos/10000^{2i}/d_{model})$
- $PE(pos, 2_{i+1}) = \cos(pos/10000^{2i}/d_{model})$



- $PE(pos, 2_i) = \sin(pos/10000^{2i}/d_{model})$
- $PE(pos, 2_{i+1}) = \cos(pos/10000^{2i}/d_{model})$



# Recap



Tokenisation breaks down natural language sequences into atomic units that carry some semantic meaning (tokens).

Encoding represents tokens in a one-dimensional numeric space (token IDs)

Embeddings project token IDs in higher dimensional vector space

Positional encoding injects information about a tokens' sequence-position into the embedding vector

*Some magic yet to be covered*

The decoder autoregressively generates a probability distribution over the vocabulary.